


Data Models

- A *data model* is a collection of concepts for describing data.
- A *schema* is a description of a particular collection of data, using the a given data model.
- The *relational model of data* is the most widely used model today.
 - Main concept: *relation*, basically a table with rows and columns.
 - Every relation has a *schema*, which describes the columns, or fields.

Levels of Abstraction

- Many views, single conceptual (logical) schema and physical schema.
 - Views describe how users see the data.
 - Conceptual schema defines logical structure
 - Physical schema describes the files and indexes used.



➡ *Schemas are defined using DDL; data is modified/queried using DML.*

Example: University Database

- Conceptual schema:
 - *Students*(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)
 - *Courses*(*cid*: string, *cname*: string, *credits*: integer)
 - *Enrolled*(*sid*: string, *cid*: string, *grade*: string)
- Physical schema:
 - Relations stored as unordered files.
 - Index on first column of Students.
- External Schema (View):
 - *Course_info*(*cid*: string, *enrollment*: integer)

Data Independence *

- Applications insulated from how data is structured and stored.
- Logical data independence: Protection from changes in *logical* structure of data.
- Physical data independence: Protection from changes in *physical* structure of data.

➡ *One of the most important benefits of using a DBMS!*

Transaction?

As services are employed for serious purposes, they will inevitably update information resources

- Can we be sure that such updates preserve integrity?
- What about when multiple services need to work together?
- What about when the services are involved in a long-lived activity?

Transactions: 1

- A transaction is a computation (i.e., program *in execution*) that accesses and possibly modifies a database:
 - Not the source code; not the binaries
 - Can be interleaved with other transactions
 - But guarantees certain correctness properties

The purpose of the transaction concept is to avoid the problems (“race conditions”) that may arise from interleaving

ACID Properties

- If programmers guarantee correctness of individual transactions, then DBMS guarantees correctness of any set of them
- The ACID properties formalize the notion of a transaction behaving as one operation
 - (Failure) *Atomicity*—all or none—if failed then no changes to DB or messages
 - This is the vernacular notion of “transaction”
 - *Consistency*—don't violate DB integrity constraints: execution of the op is correct
 - *Isolation* (Atomicity)—partial results are hidden
 - *Durability*—effects (of transactions that "happened" or committed) are forever

Transaction Lifecycle

A transaction goes through well-defined stages in its life (always *terminating*)

- Inactive
- Active (may read and write)
 - *Entire* business logic takes place here
- Precommit (no errors during execution; needed for mutual commitment protocols)
- Failed (errors)
- Committed (the DBMS decides this)
- Forgotten (the DBMS reclaims data structures)

Schedules

- Schedules are histories of computations showing all events of interest
- Schedule of $T_1 \dots T_n$ has all ops of $T_1 \dots T_n$ in the same order as within each T_i , but *interleaved* across T_i to model concurrency
 - Includes active transactions
 - Typically a partial order among events
- Two challenges
 - What are the bad schedules?
 - How can the DBMS prevent them?

Conflict

Order-sensitivity of operations

- Two operations of different transactions, but on the same data item, *conflict* if
 - Their mutual order is significant, i.e., determines at least one of the following:
 - The final value of that item read by future transactions
 - The value of the item as read by present transactions

Serial Schedules

Transactions are wholly before or after others (i.e., occur one by one)

- Clearly, we must allow for service requests to come in slowly, one-by-one
- Thus, under independence of transactions (assuming each transaction is correct), serial schedules are obviously correct

Serializable Schedules

- Interleaved schedules are desirable
 - Why?
- Those *equivalent* to some serial schedule. Here equivalent can mean
 - *Conflict* equivalent—all pairs of conflicting ops are ordered the same way
 - *View* equivalent—all users get the same view

Locks

Lock item x while using item x

- Binary: at most one party may lock x
 - Lock(x): acquire the lock
 - Computation hangs until lock(x) returns, i.e., the lock is acquired
 - Unlock(x): relinquish the lock
 - Gives mutual exclusion but restrictive

Multimode Locks

When one party has an exclusive lock, no other party may have an exclusive or a shared lock

- Shared-lock(x) needed for read(x)
 - Others can also hold a shared lock
- Exclusive-lock(x) needed for write(x)
 - No one else can concurrently hold a shared or exclusive lock
- Can be upgraded (read to write) or downgraded (write to read)

Locking

- By itself, using locks does not guarantee serializability
 - What is an example of a bad schedule obtained while using locks?
- A locking protocol, i.e., how locks are acquired and released, is critical
- That is, locks on different data items must be related

The Log

- The following actions are recorded in the log:
 - *Ti writes an object*: The old value and the new value.
 - Log record must go to disk before the changed page!
 - *Ti commits/aborts*: A log record indicating this action.
- Log records chained together by Xact id, so it's easy to undo a specific Xact (e.g., to resolve a deadlock).
- Log is often *duplexed* and *archived* on “stable” storage.
- All log related activities (and in fact, all CC related activities such as lock/unlock, dealing with deadlocks etc.) are handled transparently by the DBMS.

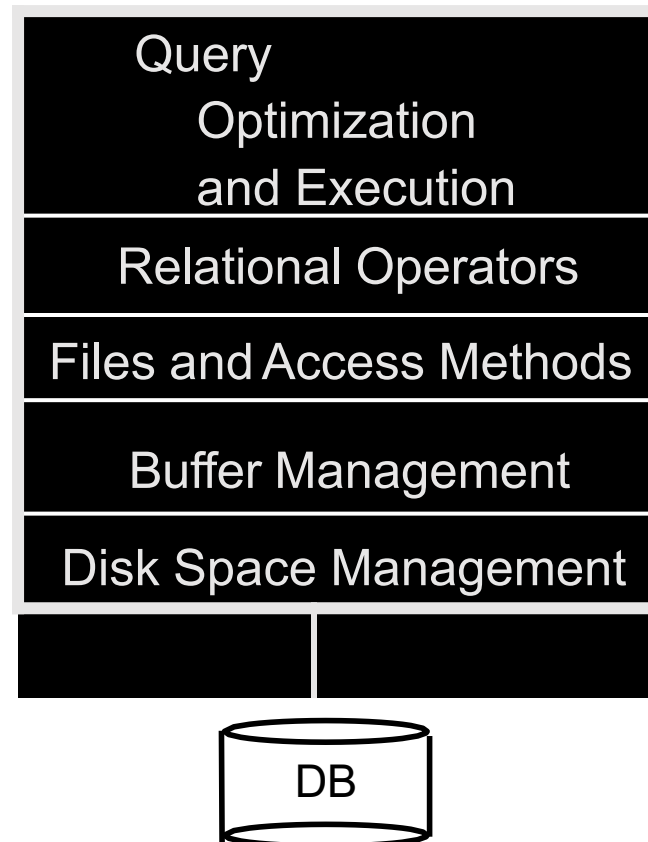
Databases make these folks happy

- End users and DBMS vendors
- DB application programmers
 - E.g., smart webmasters
- Database administrator (DBA)
 - Designs logical /physical schemas
 - Handles security and authorization
 - Data availability, crash recovery
 - Database tuning as needs evolve

Must understand how a DBMS works!

Structure of a DBMS

- A typical DBMS has a layered architecture.
- The figure does not show the concurrency control and recovery components.
- This is one of several possible architectures; each system has its own variations.



ELEMENTS OF THE DATABASE APPROACH

○ Data models

Graphical diagram capturing nature and relationship of data

Enterprise Data Model–high-level entities and relationships for the organization

Project Data Model–more detailed view, matching data structure in database or data warehouse

○ Entities

Noun form describing a person, place, object, event, or concept

Composed of attributes

○ Relationships


Between entities

Usually one-to-many (1:M) or many-to-many (M:N), but could also be one-to-one (1:1)

○ Relational Databases

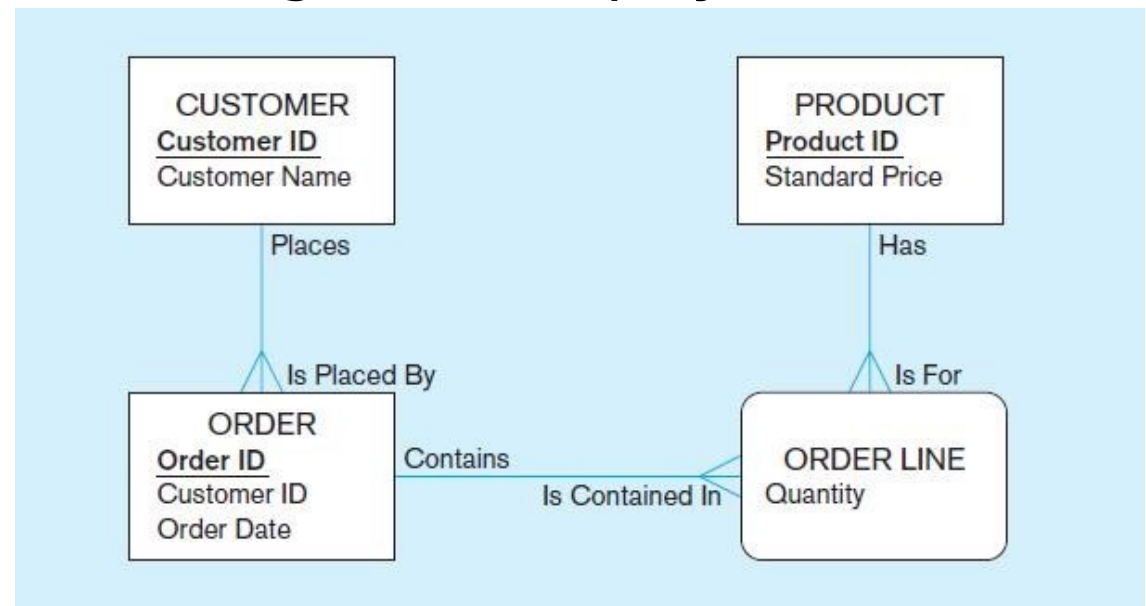
Database technology involving tables (relations) representing entities and primary/foreign keys representing relationships

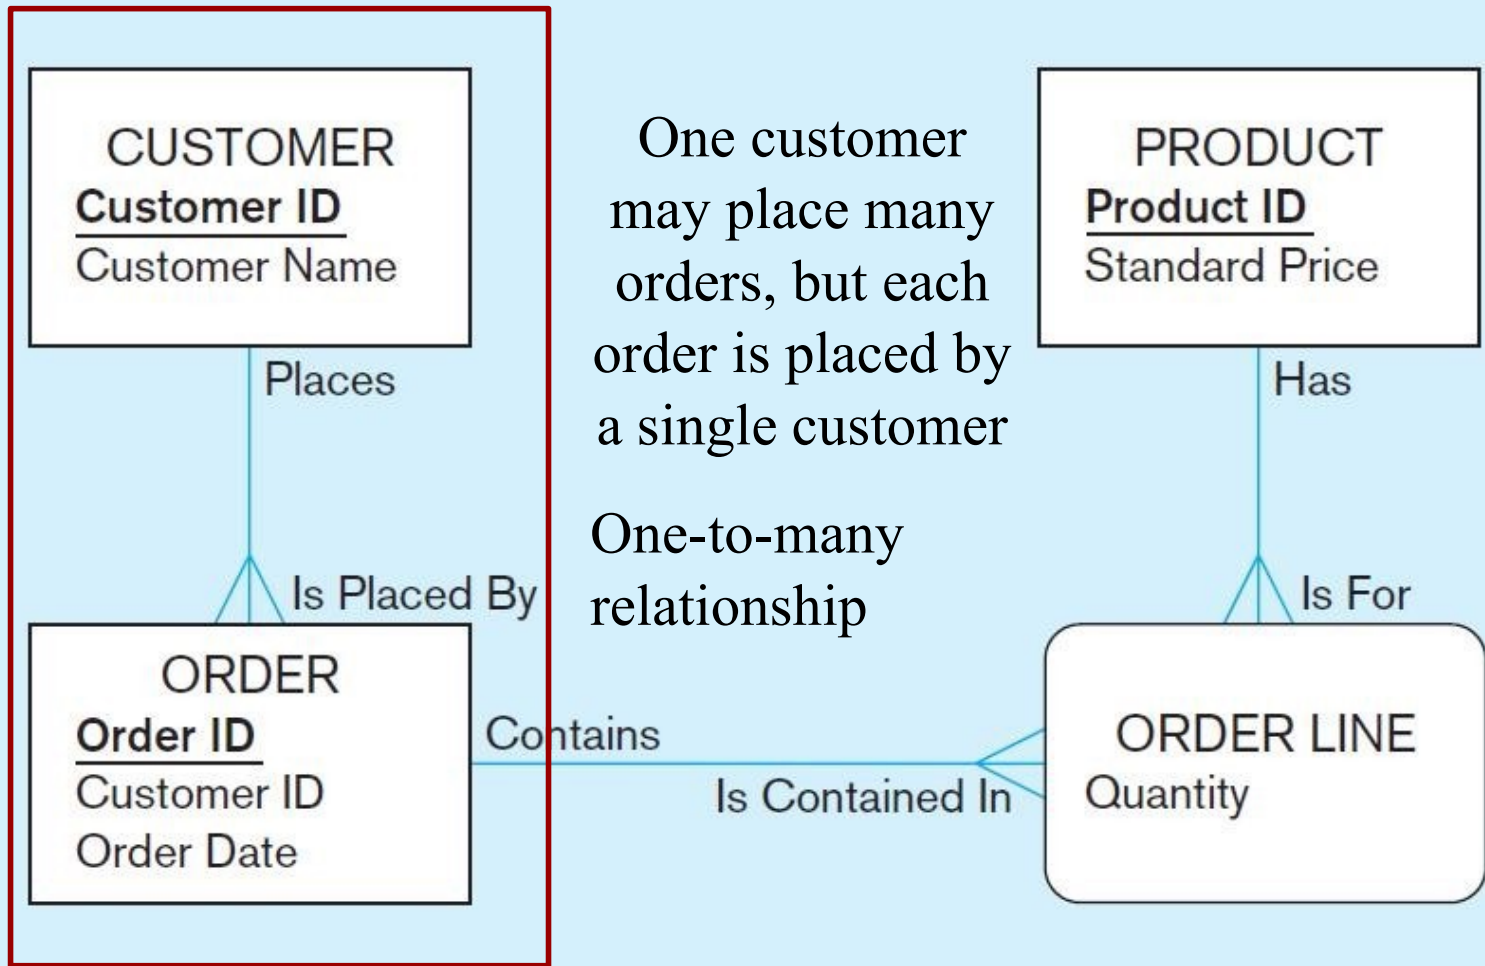
Figure 1-3 Comparison of enterprise and project level data models



Segment of an enterprise data model

Segment of a project-level





RDBMS and MySQL

Introducing Relational Databases

- A relational database manages data in tables.
- Databases are managed by a relational database management system (RDBMS).
- An RDBMS supports a database language to create and delete databases and to manage and search data.
- The database language used in almost all DBMSs is SQL.

Introducing Relational Databases

- After creating a database, the most common SQL statements used are
 - **INSERT** to add data
 - **UPDATE** to change data
 - **DELETE** to remove data
 - **SELECT** to search data
- A database table may have multiple columns, or **attributes**, each of which has a name.
- Tables usually have a **primary key**, which is one or more values that uniquely identify each row in a table

Relational Databases

Winery Table


Winery ID	Winery name	Address	Region ID
1	Moss Brothers	Smith Rd.	3
2	Hardy Brothers	Jones St.	1
3	Penfolds	Arthurton Rd.	1
4	Lindemans	Smith Ave.	2
5	Orlando	Jones St.	1

Region Table

Region ID	Region name	State
1	Barossa Valley	South Australia
2	Yarra Valley	Victoria
3	Margaret River	Western Australia

Relational Databases

- A database is modeled using **entity-relationship (ER) modeling**.



Terminology

- Database

- A repository to store data.

- Table

- The part of a database that stores the data. A table has columns or attributes, and the data stored in rows.

- Attributes

- The columns in a table. All rows in table entities have the same attributes. For example, a customer table might have the attributes name, address, and city. Each attribute has a data type such as string, integer, or date.

Terminology

- Rows

- The data entries in a table. Rows contain values for each attribute. For example, a row in a customer table might contain the values "Matthew Richardson," "Punt Road," and "Richmond." Rows are also known as records.

- Relational model

- A model that uses tables to store data and manage the relationship between tables.

- Relational database management system

- A software system that manages data in a database and is based on the relational model.

Terminology

- SQL

- A query language that interacts with a DBMS. SQL is a set of statements to manage databases, tables, and data.

- Constraints

- Restrictions or limitations on tables and attributes. For example, a wine can be produced only by one winery, an order for wine can't exist if it isn't associated with a customer, having a name attribute could be mandatory for a customer.

Terminology

- **Primary key**

- One or more attributes that contain values that uniquely identify each row. For example, a customer table might have the primary key of cust ID. The cust ID attribute is then assigned a unique value for each customer. A primary key is a constraint of most tables.

- **Index**

- A data structure used for fast access to rows in a table. An index is usually built for the primary key of each table and can then be used to quickly find a particular row. Indexes are also defined and built for other attributes when those attributes are frequently used in queries.

Terminology

- Entity-relationship modeling

- A technique used to describe the real-world data in terms of entities, attributes, and relationships.

- Normalized database

- A correctly designed database that is created from an ER model. There are different types or levels of normalization, and a third-normal form database is generally regarded as being an acceptably designed relational database.

Managing Databases

- The **Data Definition Language (DDL)** is the set of SQL statements used to manage a database.

MySQL Overview

MySQL Server

- Community Server
- Enterprise Server
- Embedded Server
- Cluster (Standard and Carrier-Grade)



MySQL GUI Tools

- Query Browser
- Administrator
- Migration Toolkit
- Visual Studio Plug-in
- MySQL Workbench



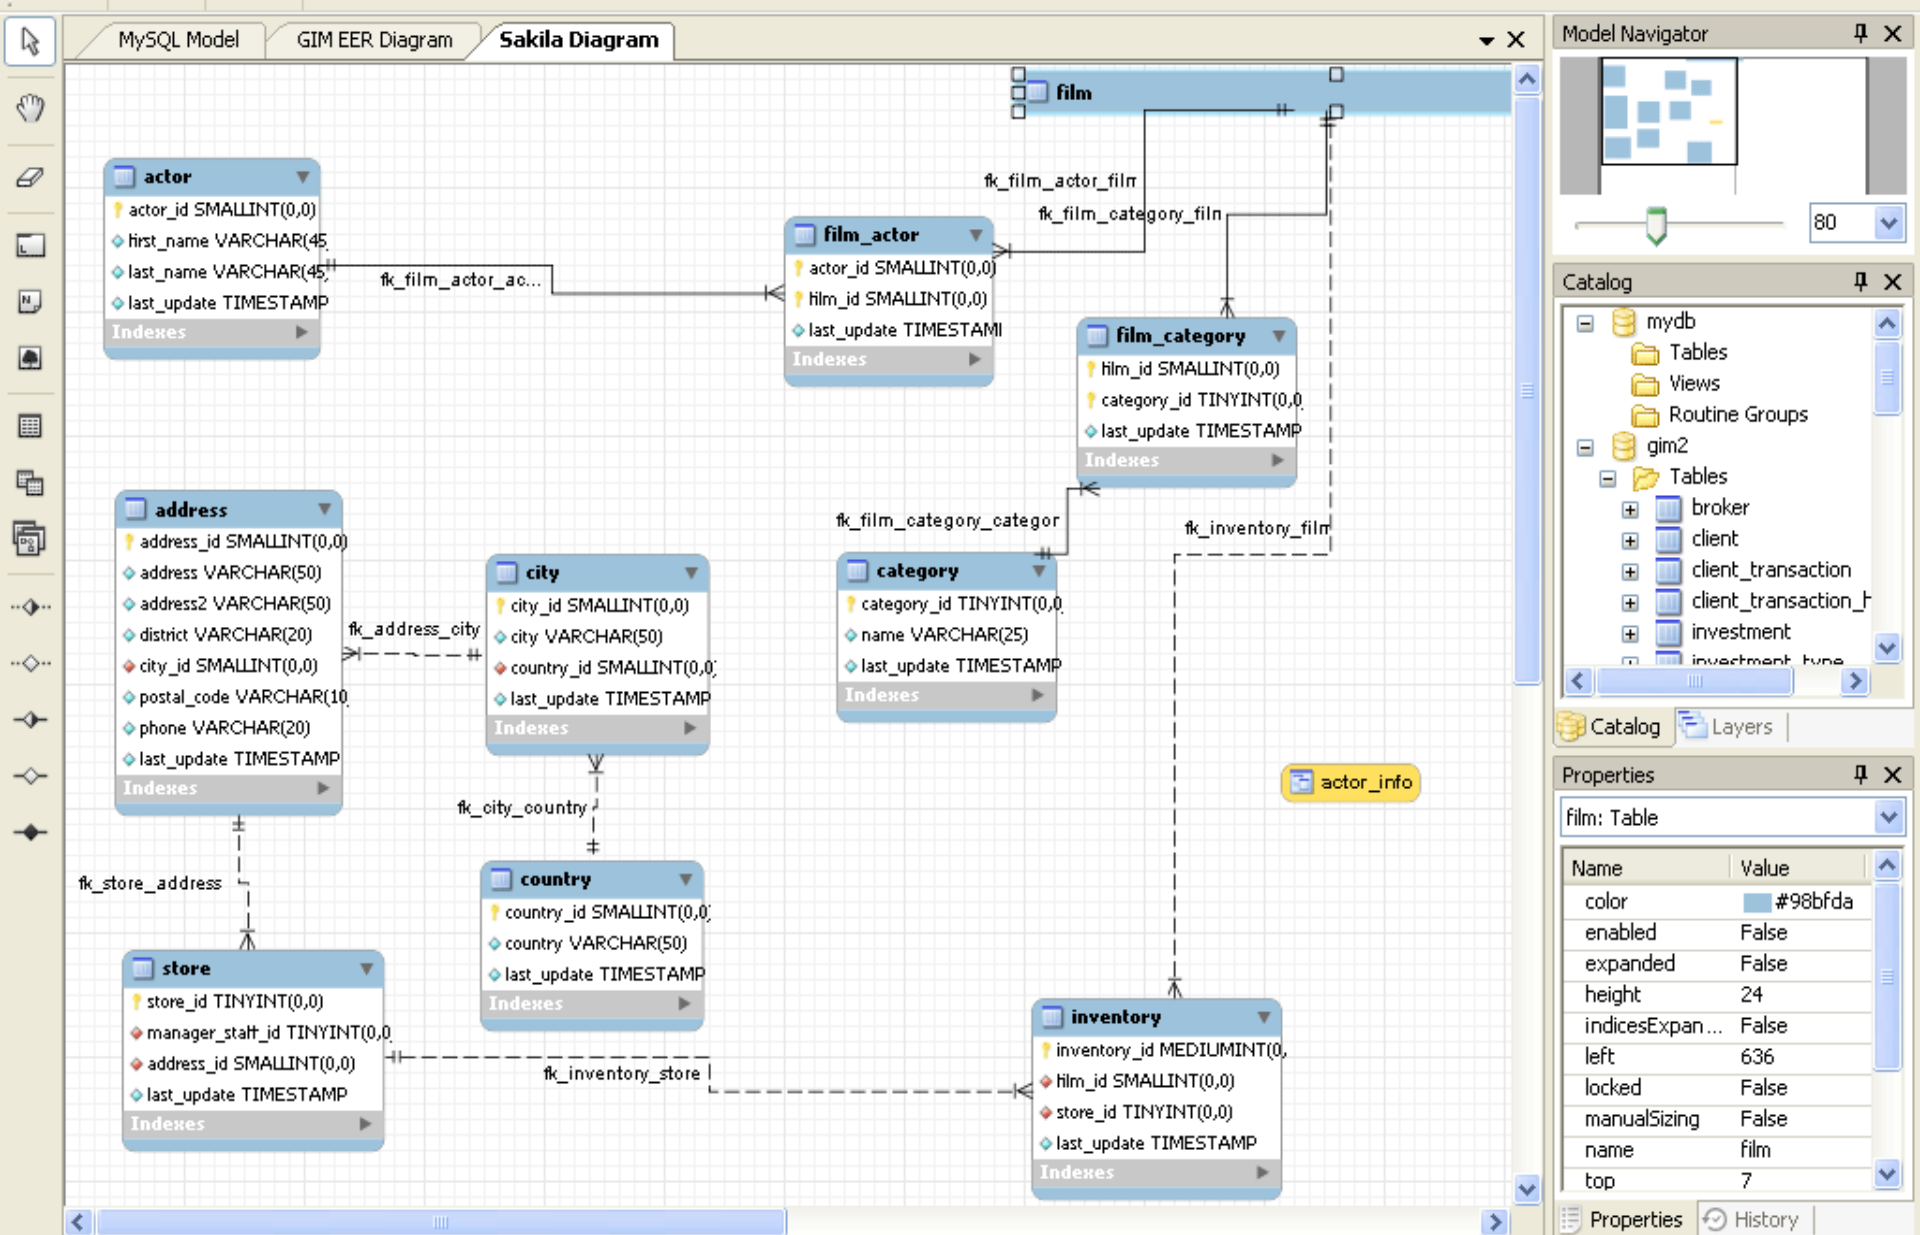
MySQL Drivers

- JDBC
- ODBC
- .NET
- PHP



MySQL Workbench

- MySQL Workbench enables a DBA, developer, or data architect to visually design, generate, and manage all types of databases including Web, OLTP, and data warehouse databases. It includes everything a data modeler needs for creating complex ER models, and also delivers key features for performing difficult change management and documentation tasks that normally require much time and effort.



Model Navigator

Catalog

- mydb
 - Tables
 - Views
 - Routine Groups
- gim2
 - Tables
 - broker
 - client
 - client_transaction
 - client_transaction_h
 - investment
 - investment_tune

Properties

Film: Table

Name	Value
color	#98bfda
enabled	False
expanded	False
height	24
indicesExpan...	False
left	636
locked	False
manualSizing	False
name	film
top	7

Properties History

MySQL Workbench

- Characteristics

❖ **Forward and Reverse Engineering**

- A visual data model can easily be transformed into a
- physical database on a target MySQL Server with just a few mouse clicks.
- it can also import SQL scripts to build models and export models to DDL scripts that can be run at a later time.

❖ **Change Management**

Sakila Sample Database