

# Introduction to NoSQL DB/MongoDB

# Advantages of RDBMS

- Relational databases are well-documented and mature technologies, and RDBMSs are sold and maintained by several established corporations.
- SQL standards are well-defined and commonly accepted.
- A large pool of qualified developers have experience with SQL and RDBMS.
- All RDBMS are ACID-compliant, meaning they satisfy the requirements of Atomicity, Consistency, Isolation, and Durability.

# Disadvantages of RDBMS

- RDBMSs don't work well with unstructured or semi-structured data due to schema and type constraints → ill-suited for large analytics or IoT event loads.
- The tables in your relational database will not necessarily map one-to-one with an object or class representing the same data.
- When migrating one RDBMS to another, schemas and types must generally be identical between source and destination tables for migration to work (schema constraint).
- For many of the same reasons, extremely complex datasets or those containing variable-length records are generally difficult to handle with an RDBMS schema.

# Types of Data

- Data can be broadly classified into four types:

## 1. Structured Data:

- Have a predefined model, which organizes data into a form that is relatively easy to store, process, retrieve and manage
- E.g., relational data

## 2. Unstructured Data:

- Opposite of structured data
- E.g., Flat binary files containing text, video or audio
- Note: data is not completely devoid of a structure (e.g., an audio file may still have an encoding structure and some metadata associated with it)

# ACID Properties in DBMS

- **Atomicity:** a transaction to be treated as a single, indivisible, logical unit of work.
- **Consistency:** a database condition in which all data integrity constraints are satisfied.
  - Consistent means that the data accurately reflects any changes up to a certain point in time.
- **Isolation:** Transactions cannot interfere with each other. Transactions are independent.
- **Durability:** Once a transaction is committed, it will remain/persist in the system, even with system failures afterwards.
  - Completed transactions persist even when servers restart or there are power failures.

# C and Latency Tradeoff

- Amazon claims that just an extra one tenth of a second on their response times will cost them 1% in sales.
- Google said they noticed that just a half a second increase in latency caused traffic to drop by a fifth.

# 4 Key Words on NoSQL

- Scale
- Speed
- Cloud
- New Data

# What is NoSQL?

- non-relational
- simple API
- schema-free
- open-source
- horizontally scalable (sharding)
- replication support
- eventually consistent /BASE



# BASE Properties in NoSQL DBMS

- The CAP theorem → impossible to guarantee strict Consistency & Availability while being able to tolerate network partitions
- NoSQL databases have relaxed ACID properties:

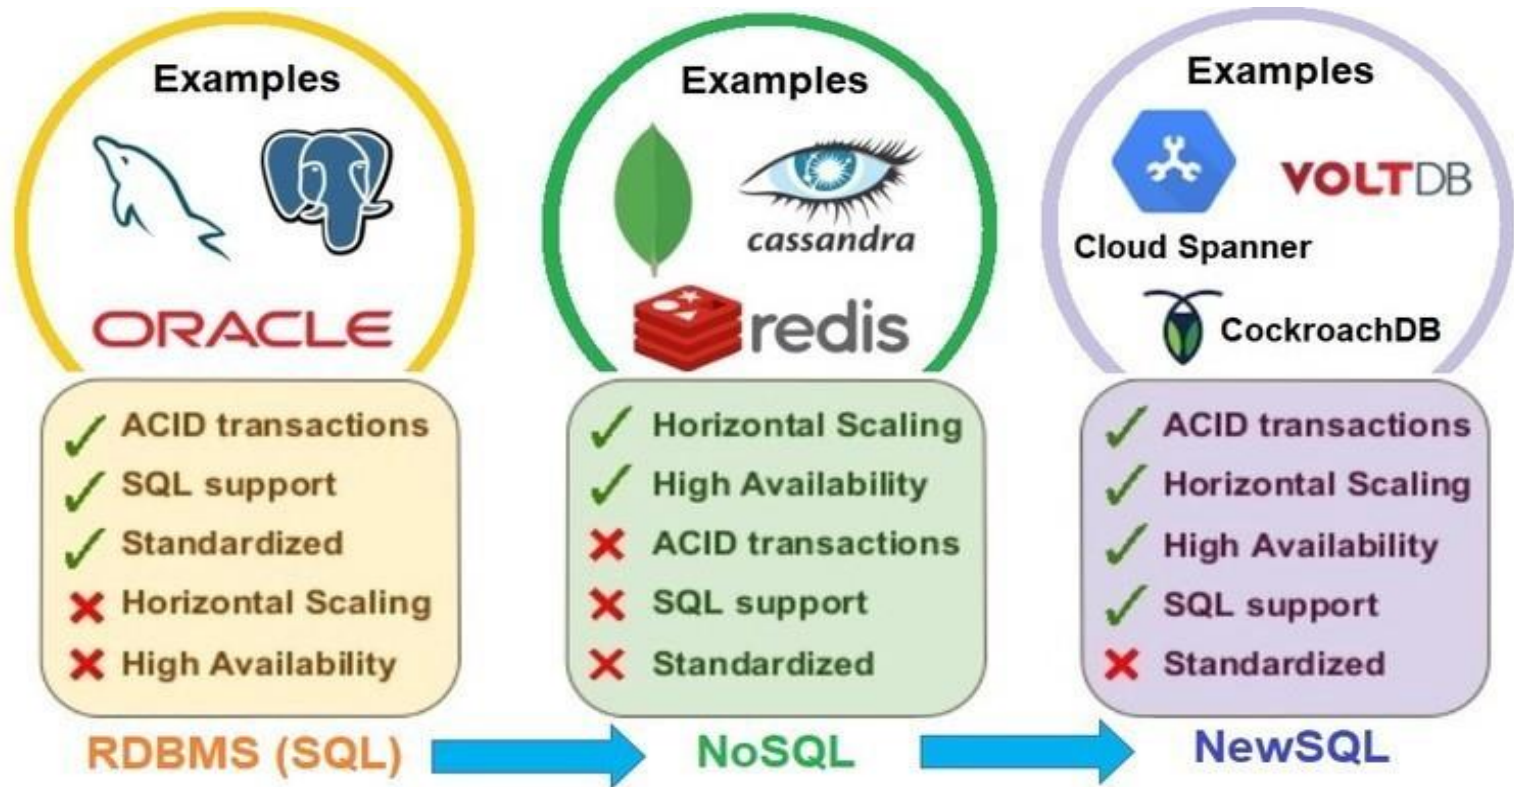
**Basically Available:** the system guarantees availability

**Soft-State:** state of the system may change over time

**Eventual Consistency:** the system will *eventually* become consistent

- Eventually Consistent means all replicas will gradually become consistent

# Examples



<https://medium.com/rabiprasadpadhy/google-spanner-a-newsql-journey-or-beginning-of-the-end-of-the-nosql-era-3785be8e5c38>

# Advantages of NoSQL DBMS

- Schema-free – allow both semi-structured and unstructured data
- Highly-scalable
- Cloud and Big data support
- Relatively light database administration

# Disadvantages of NoSQL DBMS


- Less mature technology
- Less enterprise level support
- No SQL support

# RDBMS vs. NoSQL Databases

	RDBMS	NoSQL
Variety	One type (relational)	Four main types: document, column-oriented, key-value, and graph
Structure	Predefined	Dynamic
Scaling	Primarily vertical	Primarily horizontal
Focus	Data integrity	Data performance and availability

# Types of Non-Relational Data Models/Databases

## SQL DATABASES




Relational

## NoSQL DATABASES



Column



Key-Value


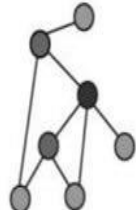

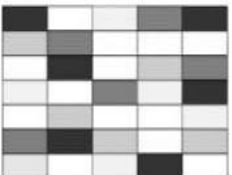
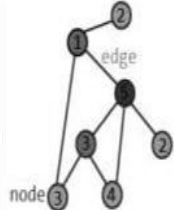

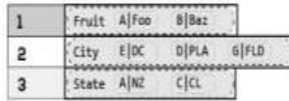






Graph



Document

# NoSQL DATABASE TYPES

Document	Graph	Key-Value	Wide-Column
			
<pre>{   "user": {     "id": "143",     "name": "improgrammer",     "city": "New York"   } }</pre>			
			

# Key-Value Stores

- Simplest type of NoSQL database
- Every database element is stored as a key-value pair
- Support CRUD operations but no join or aggregate functions
- Example: Redis, Amazon DynamoDB, Riak





# Wide-Column Stores

- Wide-column stores also referred to as column family stores
- A hybrid of RDBMS and key-value stores (aka multidimensional key-value stores):
  - Values are stored in groups/families of columns in column order (vs row order)
  - Queried by matching keys
  - Highly scalable to manage petabytes of data across massive, distributed systems
- Example: Hbase, Google BigTable, Cassandra




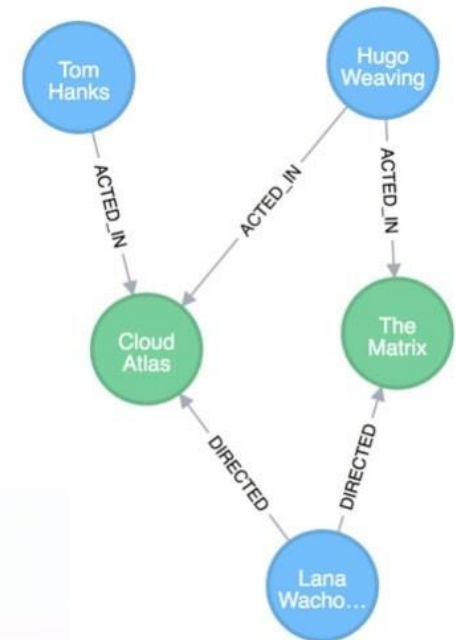

# Document Stores

- Data (document) are stored in various standard format/encoding (e.g. json, xml, pdf etc) known as BLOBS (binary large objects) although mostly JSON is used.
- Each document contains different fields including strings, dates, values and arrays.
- Documents provide intuitive and natural way to model data (hierarchical) similar to objected-oriented programming. Documents are like objects.
- Data can be indexed to outperform traditional file systems.
- Example: MongoDB, CouchDB




# Graph Stores

- Data are represented as graph structures with nodes and edges.
- Ideal for modelling complex relationships (social networking, network topologies)



# CAP Theorem

- CAP Theorem: Any distributed database with shared data can have at most 2 of the 3 properties: C, A or P



# Traditional RDMS

- Consistency
  - Transaction isolation & repeatability
  - RDMS: e.g. banking, financials,...
- Availability
  - Clustered servers
  - 99.999% uptime
- NOT easily partitioned (scalable)



# NoSQL

- Partitioning
  - Fast & global scalability
  - NoSQL: big data, mobile gaming, social media
- Availability
  - Clustered servers
  - 99.999% uptime
- Sacrifice consistency

# NoSQL Database: Major Players

- Too many document NoSQL databases to name a few distinct ones

29 systems in ranking, July 2014

Rank	Last Month	DBMS	Database Model	Score	Changes
1.	1.	<a href="#">MongoDB</a>	Document store	238.78	+7.33
2.	2.	<a href="#">CouchDB</a>	Document store	23.07	+0.28
3.	3.	<a href="#">Couchbase</a>	Document store	16.58	+0.79
4.	4.	<a href="#">MarkLogic</a>	Multi-model 	8.20	-0.02
5.	5.	<a href="#">RavenDB</a>	Document store	5.09	-0.42
6.	6.	<a href="#">GemFire</a>	Document store	2.16	-0.06
7.	7.	<a href="#">OrientDB</a>	Multi-model 	1.71	-0.02
8.	8.	<a href="#">Cloudant</a>	Document store	1.70	+0.07
9.	9.	<a href="#">Datameer</a>	Document store	0.88	+0.08
10.	10.	<a href="#">Mnesia</a>	Document store	0.72	+0.01

# Key Benefit of NoSQL: $O(1)$ Lookup

- Fast lookup
  - No joining required
  - All data about one domain concept in one document
- Direct programming language representation
  - No mapping or ‘ORM’ layer required
- JSON library
  - Direct result representation and manipulation
  - JavaScript: representation in language data types directly
  - E.g., check out MongoDB node.js driver



# MongoDB

<https://en.wikipedia.org/wiki/MongoDB>

- Name derived from Hu(**MONGO**)us word
- Document Oriented Database
- Built for High – Performance and scalability
- Document based queries for **Easy Readability**
- Replication and failover for **High Availability**
- Auto Sharding for **Easy Scalability**

# Where to use MongoDB ?

- Ideal for Web Applications
- Applications containing semi-structured data and need flexible schema management
- Caching and High Scalability
- Scenarios where **data availability** and **size of data** are priorities over the **transactions** of data

# How does MongoDB Store data?

- Stores data in form of Documents
- JSON like field – value pair
- Documents analogous to structures in programming languages with key – value pair
- Documents stored in **BSON (Binary JSON)** format
- BSON is JSON with additional type information

# Companies Using MongoDB



"MongoDB gives you the ability to concentrate on your business and create the applications. Everything else is taken care of."  
— Steven Bondi, Forbes.com



"The flexibility of MongoDB has put a lot of power in the hands of our developers."  
— Brian Massey, Under Armour



"Personalization based on real-time data is the key success factor for e-commerce sites."  
— Peter Wolter, OTTO



"I want to be the city that's on the forefront, actually building a new path, building a new model for how cities operate."  
— City of Chicago



"As part of our infrastructure redesign, we needed to ensure that new app development was never waiting on the backend."  
— Luke Kolin, The Weather Channel




"The best thing from the get-go was MongoDB's engineers and developer-first approach."  
— Tony Tan, Reverb Technologies



# MongoDB: A Document Store

- A record in MongoDB is a document, which is a data structure composed of field and value pairs.
- MongoDB documents are similar to JSON objects. The values of fields may include other documents, arrays, and arrays of documents.

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```



← field: value  
← field: value  
← field: value  
← field: value

# MongoDB Terminology

- Mysql
  - Table
  - Row
  - Column
  - Joins
  - Group By
- MongoDB

# MongoDB Terminology

- Mysql
- Table
- Row
- Column
- Joins
- Group By
- MongoDB
- Collection
- Document
- Field
- Embedded document
- Aggregation

# Examples of MongoDB Data

```
{
  "_id": {
    "$oid": "56d61033a378eccde8a8357e"
  },
  "address": {
    "city": "LAWRENCE",
    "number": 1,
    "street": "BAY BLVD",
    "zip": 11559
  },
  "business_name": "SPRAGUE OPERATING RESOURCES LLC.",
  "certificate_number": 3019422,
  "date": "Mar 3 2015",
  "id": "11247-2015-ENFO",
  "result": "Fail",
  "sector": "Fuel Oil Dealer - 814"
}
```

```
{
  "address": {
    "building": "8825",
    "coord": [-73.8803827, 40.7643124],
    "street": "Astoria Boulevard",
    "zipcode": "11369"
  },
  "borough": "Queens",
  "cuisine": "American",
  "grades": [ {
    "date": {"$date": "2014-11-15T00:00:00.000Z"},
    "grade": "Z",
    "score": 38
  },
  {
    "date": {"$date": "2014-05-02T00:00:00.000Z"},
    "grade": "A",
    "score": 10
  },
  {
    "date": {"$date": "2013-03-02T00:00:00.000Z"},
    "grade": "A",
    "score": 7
  },
  {
    "date": {"$date": "2012-02-10T00:00:00.000Z"},
    "grade": "A",
    "score": 13
  }
],
  "name": "Brunos On The Boulevard",
  "restaurant_id": "40356151"
}
```



# SQL vs MongoDB



relational database tables in a 1-many relationship

```
{
  first_name: "Paul",
  surname: "Miller",
  city: "London",
  location: [45.123,47.232],
  cars: [
    { model: "Bentley",
      year: 1973,
      value: 100000, ...},
    { model: "Rolls Royce",
      year: 1965,
      value: 330000, ...},
  ]
}
```

a single document (collection)

# Document in MongoDB

- Stored in Collections
- Has **\_id** field – works like Primary keys in Relational databases
- Sample document containing name, age, status and groups

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value  
← field: value  
← field: value  
← field: value

# MongoDB Query Language (MQL)

- Field selection
- The find () function in MongoDB Query Language provides the easiest way to retrieve data from multiple documents within one of your collections.
- `mql>db.Employee.find()`
- `Sql>select * from Employee;`

# **MQL:** Condition clause: Where clause

- Here, to return all the A-status employees
- `mql>db.Employee.find ( { status : "A" } )`
- `sql>select * from Employee where status='A'`

# **MQL:** Projection: Select specific fields

- The insertion of the tag { NAME : 1 } specifies that only the information from the NAME field should be returned. The results are sorted and presented in ascending order.
- `mql>db.Employee.find ( { status : "A" }, {NAME: 1 }  
);`
- `sql>select name from Employee where status='A';`

# **MQL:** Field exclusion

- To query for the opposite, inserting { NAME : 0 } retrieves a list of all fields except for the NAME field.
- `mql>db.Employee.find ( {status : "A"}, {NAME: 0} );`
- `sql>select age, address, status from Employee;`

# **MQL:** Complex operation (Dot notation)

- When you work with more complex document structures such as documents containing arrays or embedded objects, you can use other methods for querying from them.
- `mql>db.Employee.find( { "Address.City" : "Irving" }).pretty()`
- `sql>SELECT E.* from Employee E inner join Dept D on E.EMPID=D.EMPID Where D.City='Irving'`

# MQL: Insert multiple documents

- To create an array of documents, define the variable by a name and assign the array of documents.
- `mql>document = [ { "Name" : "Brian Lockwood", "Age" : "45", status:"A"}, { "Name" : "Charles", "Age" : "35", status:"A"} ]`  
`>db.Employee.insert (document)`
- `sql>Insert into Employee values ('Brian Lockwood',45, 'A'),('Charles',35,'A')`



# **MQL:** Insert nested document

- the address document is embedded in the document. MongoDB Query Language accommodates that, but SQL strictly follows procedural constructs and does not allow insertion of values to non-existing fields.
- `mql>document = ({ "Name" : "Robert Jordan","Age" : "37", status:"A", Address: { Street:"Polaris Way", City : "Aliso Viejo",State:"California" } })`
- `>db.Employee.insert (document)`
- `sql>Not applicable`

# **MQL:** Get distinct status

- To return unique values only
- `mql>db.Employee.distinct( "status");`
- `sql>Select distinct status from Employee;`

# MQL: Sort

- This example sorts the results based on the Age key in ascending order. Sorting is in ascending order unless otherwise specified (-1 flag for descending order).
- `mql>db.Employee.find().sort( { Age: 1 } )`
- `>db.Employee.find().sort( { Age: -1 } )`
- `sql>Select * from Employee order by Age ASC;`
- `Select * from Employee order by Age DESC;`

# MQL: Limit

- Use the limit () function in MongoDB Query Language to specify the maximum desired number of results
- `mql>db.Employee.find().limit(5);`
- `>db.Employee.skip().limit(5);`
- `sql>select top(5) * from Employee;`
- `SELECT * FROM Employee OFFSET 5 ROWS;`

# MQL: Aggregate

- `count ()` returns the number of documents in the specified collection.
- `mql>db.Employee.find().count();`
- `sql>Select count(*) from Employee;`

# MQL: Group

- `group()` takes three parameters: key, initial, and reduce. The purpose of `group ()` and SQL's `GROUP BY` is to return an array of grouped items.
- `mql>db.Employee.aggregate([ { "$group" : { _id:"$status",  
count:{ $sum:1 } } } ])`
- `sql>Select status,count(*) from Employee Group by status`

# MQL: Comparison

- Use special parameters \$gt, \$lt, \$gte and \$lte to perform greater-than and less-than comparisons in queries.
- `mql>db.Employee.find ( { Age: { $gt : "30" } });`
- `sql>Select * from Employee where Age>30`

# **MQL:** Multiple expressions in document

- Use \$or to search for multiple expressions in a single query.
- `mql>db.Employee.find({ $or : [ { "Name" : "Flynn" }, { "status" : "A" } ] })`
- `sql>select * from Employee where name='Flynn' or Status='A'`



# MQL: Comparison

- Use special parameters \$gt, \$lt, \$gte and \$lte to perform greater-than and less-than comparisons in queries.
- `mql>db.Employee.find ( { Age: { $gt : "30" } });`
- `sql>Select * from Employee where Age>30`

# MQL

- Combine multiple operators
- Can include limits, skips and sorts

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection  
← query criteria  
← projection  
← cursor modifier

# MongoDB Example: Create Table

- No need to create tables (collections) in MongoDB → just start inserting data

## SQL

```
CREATE TABLE people (  
  id MEDIUMINT NOT NULL  
    AUTO_INCREMENT,  
  user_id Varchar(30),  
  age Number,  
  status char(1),  
  PRIMARY KEY (id)  
)
```

## MongoDB

```
db.people.insertOne( {  
  user_id: "abc123",  
  age: 55,  
  status: "A"  
} )
```

However, you can also explicitly create a collection:

```
db.createCollection("people")
```

# MongoDB Example: Update Data

- Use updateMany() to update records

## SQL

```
UPDATE people  
SET status = "C"  
WHERE age > 25
```

```
UPDATE people  
SET age = age + 3  
WHERE status = "A"
```

## MongoDB

```
db.people.updateMany(  
  { age: { $gt: 25 } },  
  { $set: { status: "C" } }  
)
```

```
db.people.updateMany(  
  { status: "A" } ,  
  { $inc: { age: 3 } }  
)
```

# MongoDB Example: Delete Data

- Use deleteOne() or deleteMany() to delete records

## SQL

```
DELETE FROM people  
WHERE status = "D"
```

```
DELETE FROM people
```

## MongoDB

```
db.people.deleteMany( { status: "D" } )
```

```
db.people.deleteMany({})
```

# Remove Operation

- In MongoDB, `db.collection.remove()` method deletes document from the collection

```
db.users.remove(  
    { status: "D" }  
)
```

← collection

← remove criteria

# MongoDB Example: Alter Table

- Collections do not describe or enforce the structure of its documents

## SQL

```
ALTER TABLE people  
ADD join_date DATETIME
```

```
ALTER TABLE people  
DROP COLUMN join_date
```

## MongoDB

```
db.people.updateMany(  
  { },  
  { $set: { join_date: new Date() } }  
)
```

```
db.people.updateMany(  
  { },  
  { $unset: { "join_date": "" } }  
)
```

# MongoDB Example: Create Index & Drop Table

- MongoDB provides a `createIndex()` method to **create one or more indexes on collections**

```
CREATE INDEX idx_user_id_asc  
ON people(user_id)
```

```
db.people.createIndex( { user_id: 1 } )
```

```
CREATE INDEX  
    idx_user_id_asc_age_desc  
ON people(user_id, age DESC)
```

```
db.people.createIndex( { user_id: 1, age: -1 } )
```

```
DROP TABLE people
```

```
db.people.drop()
```



# What is Mongoose/**Mongothon**

- Mongoose is an [Object Document Mapper](#) (ODM) for Node.js that makes using MongoDB easier by translating documents in a MongoDB database to objects in the program.
- This means that Mongoose allows you to define objects with a strongly-typed schema that is mapped to a MongoDB document.
- **Mongothon** is a MongoDB object-document mapping API for Python, loosely based on the awesome mongoose.
- [Doctrine](#) (library that provides a PHP object mapping),
- [MongoLink](#) (object/document mapper for java)
- [Mandango](#) (The easy, powerful and ultrafast ODM for PHP).

# Mongoose

Schema Types that a property is saved as when it is persisted to MongoDB.

- String
- Number
- Date
- Buffer
- Boolean
- Mixed
- ObjectId
- Array

Each data type allows you to specify:

- a default value
- a custom validation function
- indicate a field is required
- a get function that allows you to manipulate the data before it is returned as an object
- a set function that allows you to manipulate the data before it is saved to the database
- create indexes to allow data to be fetched faster