

***Google Big Query  
and  
ElasticSearch***

# Relational Databases at Scale?

- How do traditional relational databases handle record growth at scale?
- Traditionally, very large tables are hard to scan & compute

Organization Details

Company ID	Company Name
161218560	NY Association Inc.
...	...
...	...
...	...
10 Billion Row Table	

# Relational Databases at Scale?

- Indexes (pre-sorted) help with common queries:


Organization Details

Company ID	Company Name
161218560	NY Association Inc.
...	...
...	...
...	...
10 Billion Row Table	

Index


Company Name	Ranked Order
ACME Inc.	1
...	...
...	
NY Association Inc.	900,000
...	

# 10+ Years of Tackling Big Data Problems



# What is Google BigQuery?

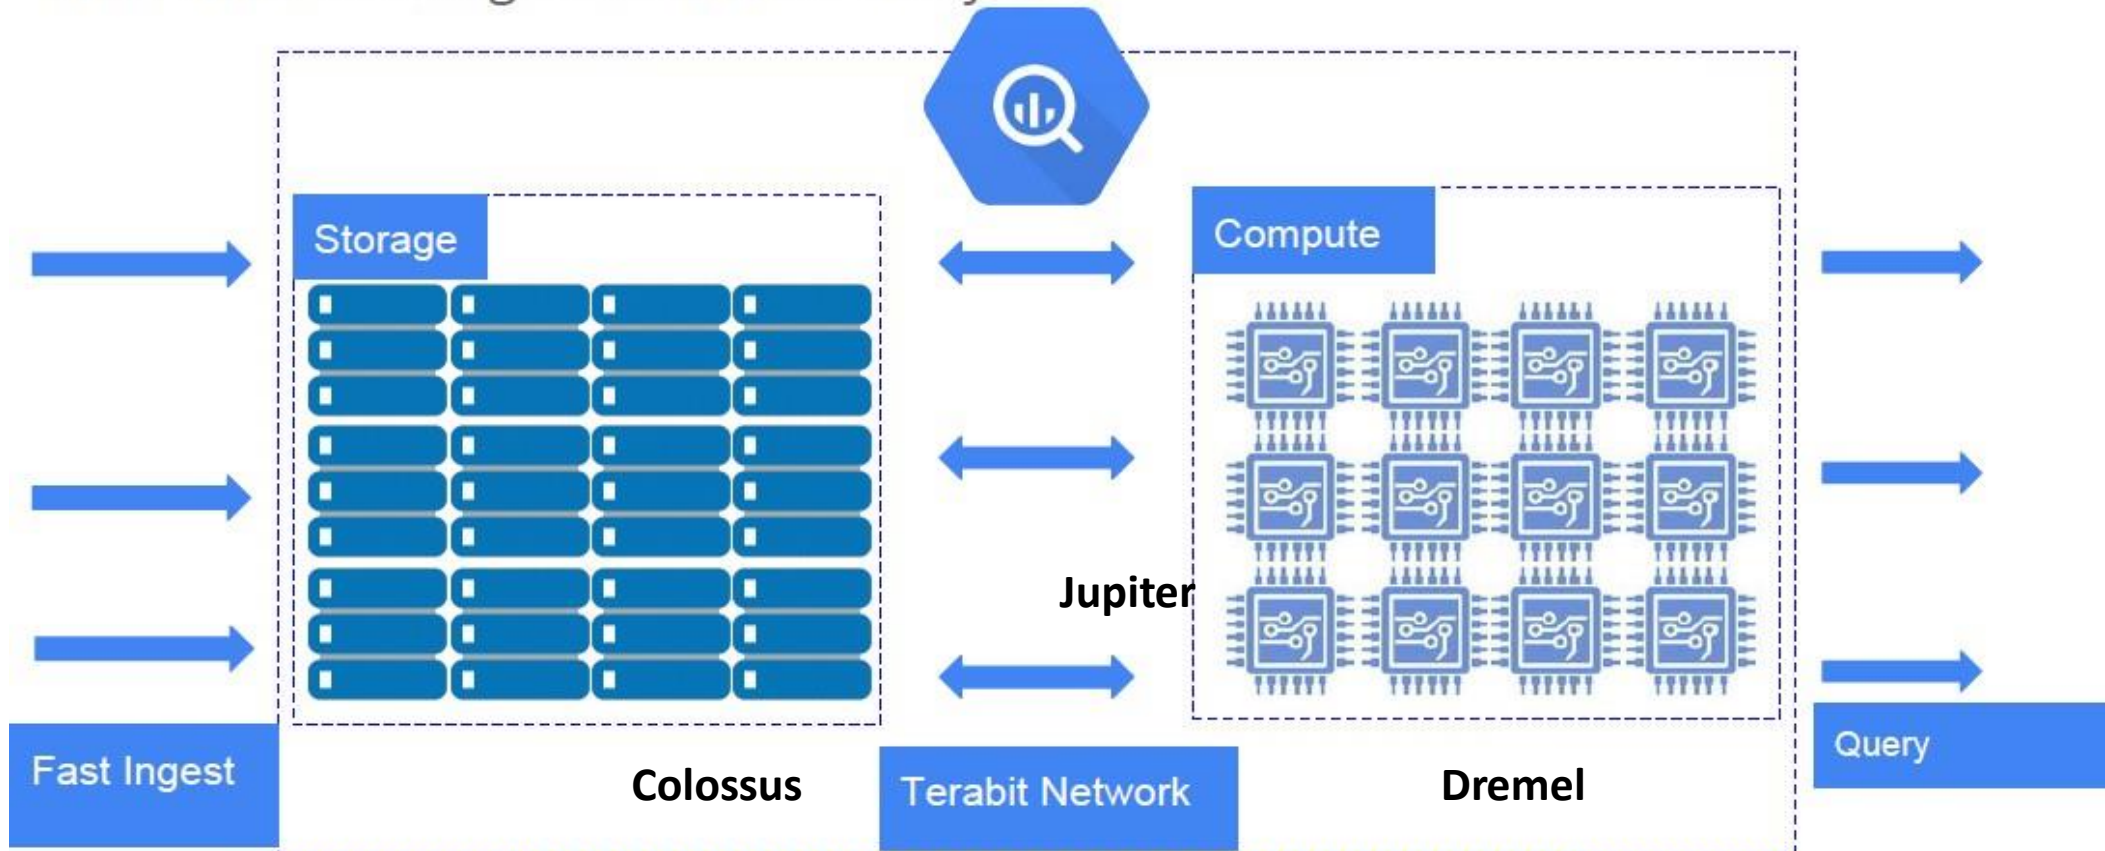
- BigQuery is a **hybrid** system that allows you to store data in columns
- A serverless, highly scalable, and cost-effective multi-cloud data warehouse designed for business agility
- A “hybrid SQL-NoSQL” database:
  - **NoSQL** (JSON like) storage (under the hood)
  - **SQL** query interface
- Extended to support machine learning
- Ideal for big data application



# A Distributed System for Querying Very Large Datasets

Google BigQuery

The Power of Google Dremel for everyone



Dremel turns SQL queries into execution trees



# Google BigQuery Performance

100B Benchmark with 3 wildcards ? Query Editor UDF Editor X

```
1 SELECT language, SUM(views) as views
2 FROM [bigquery-samples:wikipedia_benchmark.Wiki100B]
3 WHERE REGEXP_MATCH(title, "G.*o.*o.*g")
4 GROUP BY language
5 ORDER BY views desc;
```

No Cached Results X

RUN QUERY Save Query Save View Format Query Show Options

Query complete (24.7s elapsed, 4.06 TB processed) ✓

Running an *inefficient* regular expression over 100 billion rows in  
**less than 60 seconds**



**Big Data at Scale!!!**

# Why is BigQuery so fast?

**BigQuery** is designed to query structured and semi-structured data **using standard SQL**.

It is highly optimized for query performance and provides extremely high cost effectiveness.

BigQuery is a cloud-based fully-managed service

## **Columnar Storage:**

BigQuery stores data in a proprietary columnar format called Capacitor. Data is stored in a columnar storage fashion which makes possible to achieve a very **high compression** ratio and **scan throughput**.

Tree Architecture is used for dispatching queries and aggregating results across thousands of machines in a few seconds.



# Google BigQuery vs Relational Databases

	Relational DBMS	BigQuery
SQL queries	Yes	Yes
Data	Structured	Structured + <b>Nested Data</b>
Index-based?	Index	Non-Index
Storage	Row based	Column based
ACID Transactions?	ACID	<b>No</b>

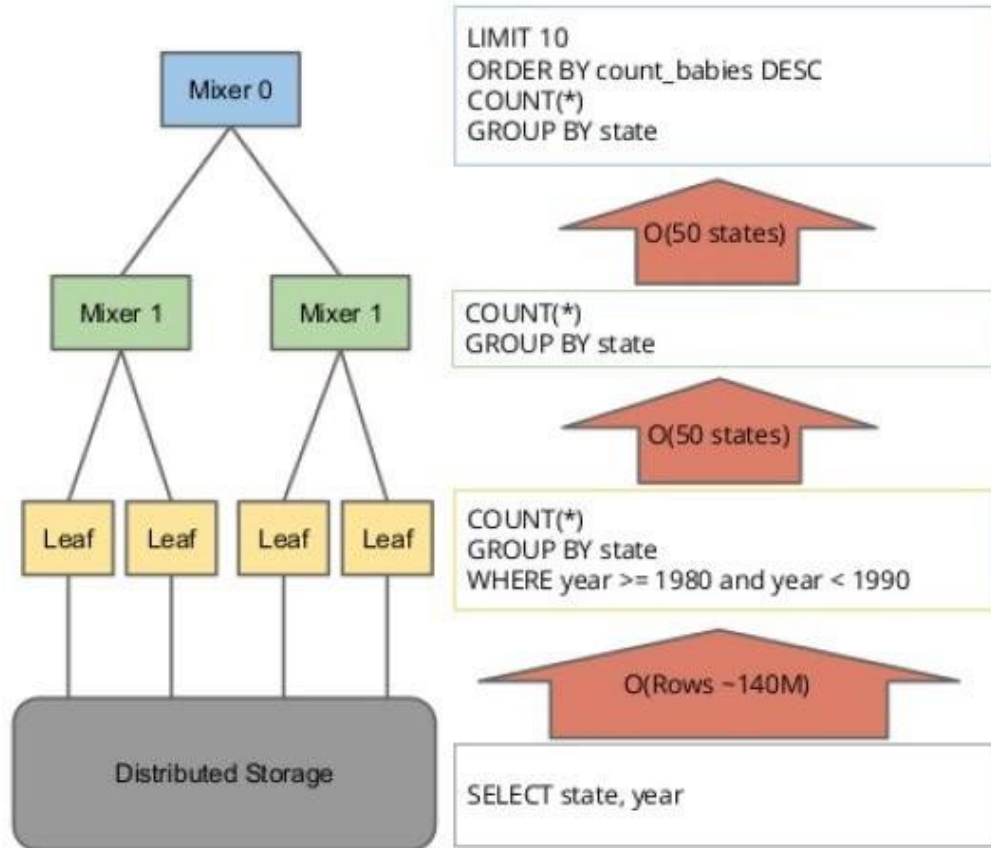
<https://db-engines.com/en/system/Google+BigQuery%3BMySQL>

# Google BigQuery

- Google BigQuery (GBQ) introduces several key innovations for scalability
- **Column-based** data storage
- **Break Apart Tables** into pieces
- **Nested Fields** within a table

# How BigQuery works

## Tree Structured Query Dispatch and Aggregation




```
SELECT
  state, COUNT(*) count_babies
FROM [publicdata:samples.natality]
WHERE
  year >= 1980 AND year < 1990
GROUP BY state
ORDER BY count_babies DESC
LIMIT 10
```

# BigQuery Column-Oriented Storage



Record Oriented Storage

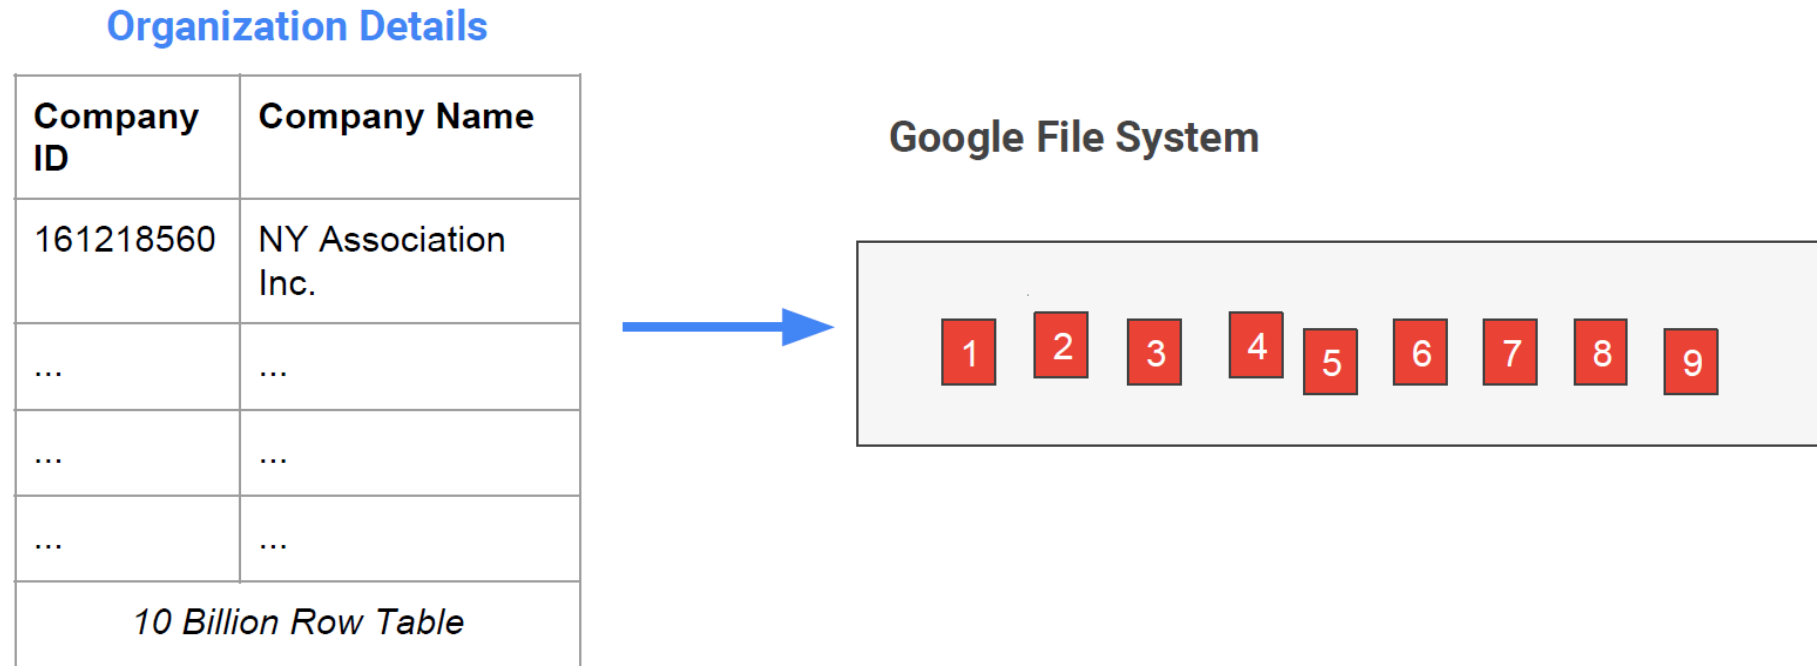


Column Oriented Storage

- Storing related values (faster to loop through at execution time)
- Columns can be individually compressed
- Access values from a few columns without reading every one

# BigQuery Data Sharding

- BigQuery automatically breaks apart data into smaller shards



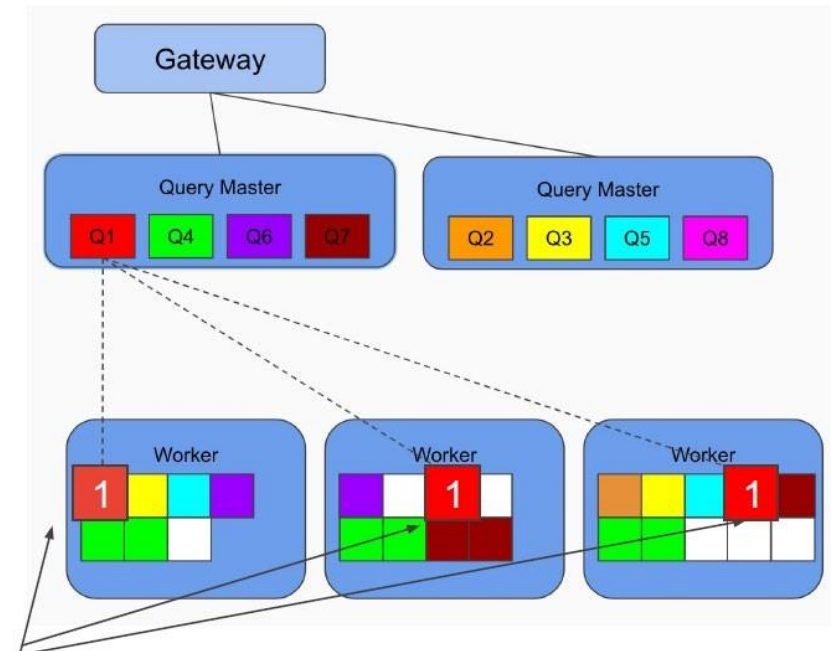
# BigQuery Data Sharding

- BigQuery automatically pieces it back together for queries

Organization Details

Company ID	Company Name
161218560	NY Association Inc.
...	...
...	...
...	...
10 Billion Row Table	

SELECT Company Name ORDER BY Company Name



Shards of data are read and processed in **parallel**



# BigQuery: Normalization vs Denormalization vs Nested Data

Normalized

people	cities_lived
<b>name</b> age gender	<b>name</b> city years_lived

Denormalized

people_cities_lived
<b>name</b> age gender city_name years_lived

Repeated

people_cities_lived
<b>name</b> age gender cities_lived ( <i>repeated</i> ) city years_lived

Less Performant

High Performing

# BigQuery Nested & Repeated Fields

- BigQuery can use **nested schemas** for highly scalable queries

Organization Details with Nested Historical Transactions

NESTED

Company ID	Company Name	Transactions.Amount	Code.Expense
161218560	NY Association Inc.	\$10.000	Lobbying
		\$5,000	Legal
		\$1,000	Insurance
123435560	ACME Co.	\$7,000	Travel


- Avoid costly joins
- No performance slowdown for `SELECT(DISTINCT Company ID)`

# Arrays and Structs in BigQuery

- Nested and repeated data are extensions to Standard SQL (SQL 2011 standard) so their supports vary among vendors.
- <https://cloud.google.com/bigquery/docs/reference/standard-sql/arrays>
- ARRAYS and STRUCTs are powerful concepts
- Nested & repeated records are **ARRAYs of STRUCTs**
- Tables with repeated fields are conceptually like pre-joined tables

# Arrays and Structs in BigQuery

- **ARRAYs** are ordered lists of zero or more data values that have the same data type:



- **STRUCTs** are flexible containers of ordered fields each with a type (required) and field name (optional)
  - Can store multiple data types (including ARRAYs) in a STRUCT



[BigQuery Nested and Repeated Fields: Dig Deeper into Data \(Cloud Next '18\)](#)



- Nested records in BigQuery are Arrays of Structs.
- Instead of Joining with a `sql_on:` expression, **the join relationship is built into the table.**
- UNNESTing a ARRAY of STRUCTs is similar to joining a table.

<https://cloud.google.com/bigquery/docs/reference/standard-sql/arrays>

<https://cloud.google.com/bigquery/docs/reference/standard-sql/dml-syntax>

## Array Example

build an array literal in BigQuery using brackets ([ and ]). Each element in an array is separated by a comma. **SELECT** [1, 2, 3] **as** numbers; **SELECT** ["apple", "pear", "orange"] **as** fruit;

```
SELECT  
['raspberry', 'blackberry', 'strawberry', 'cherry']  
AS fruit_array
```



Row	fruit_array
1	raspberry
	blackberry
	strawberry
	cherry

```
WITH fruits AS (  
  SELECT ['raspberry', 'blackberry', 'strawberry', 'cherry']  
  AS fruit_array  
)  
  
SELECT ARRAY_LENGTH(fruit_array) AS array_size  
FROM fruits;
```



Row	array_size
1	4

<https://cloud.google.com/bigquery/docs/reference/standard-sql/arrays>



## STRUCT Example

```
SELECT  
STRUCT(35 AS age, 'Jacob' AS name)
```



Row	f0_.age	f0_.name
1	35	Jacob

```
SELECT  
STRUCT(35 AS age, 'Jacob' AS name) AS customers
```



Row	customers.age	customers.name
1	35	Jacob

# STRUCT & ARRAY Example

<https://cloud.google.com/bigquery/docs/nested-repeated>

- STRUCTs Can Even Contain ARRAY Values

```
SELECT  
STRUCT(35 AS age, 'Jacob' AS name, ['apple', 'pear', 'peach'] AS  
items) AS customers
```



Row	customers.age	customers.name	customers.items
1	35	Jacob	apple
			pear
			peach

# STRUCT & ARRAY Example

- ARRAYS can Contain STRUCTs as Values

```
SELECT  
[  
  STRUCT(35 AS age, 'Jacob' AS name, ['apple', 'pear', 'peach'] AS items),  
  STRUCT(33 AS age, 'Miranda' AS name, ['water', 'pineapple', 'ice cream'] AS items)  
] AS customers
```



Row	customers.age	customers.name	customers.items
1	35	Jacob	apple
			pear
			peach
	33	Miranda	water
			pineapple
			ice cream

# Example: Data Warehouse Schema

Original Data

OrderId	CustomerId	CustomerName	timestamp	Location	purchasedItems			
					sku	description	quantity	price
1000001	65401	John Doe	12/18/2017 15:00	Faraway	ABC123456	Redwood 8x4	3	36.3
					TBL535522	Sapient Table	1	878.4
					CHR762222	Cherrywood Ch	6	435.6
					sku	description	quantity	price
1000002	74682	Jane Michaels	12/16/2017 11:30	Nearland	GCH635354	Garden chairs	4	345.7
					GRD828822	Ceramic Pots	2	9.5
					sku	description	quantity	price
1000003	63636	Jose Carlos	12/16/2017 13:40	Nearland				




Transaction Fact					
Order Id	timestamp	CustomerId	sku	quantity	price
1000001	12/18/2017 15:02:00	65401	ABC123456	3	36.3
1000001	12/18/2017 15:02:00	65401	TBL535522	1	878.4
1000001	12/18/2017 15:02:00	65401	CHR762222	6	435.6
1000002	12/16/2017 11:34:00	74682	GCH635354	4	345.7
1000002	12/16/2017 11:34:00	74682	GRD828822	2	9.5

Customer Dimension		
CustomerId	CustomerName	Location
65401	John Doe	Faraway
74682	Jane Michaels	Nearland
63636	Jose Carlos	Nearland

Product Dimension	
sku	description
ABC123456	Redwood 8x4
TBL535522	Sapient Table
CHR762222	Cherrywood Chair
GCH635354	Garden chairs
GRD828822	Ceramic Pots


# Example: Data Warehouse in BigQuery (Denormalizing data with nested and repeated structures)

## Nested Fields




- **STRUCT/RECORD** data type contains ordered fields with a type and name.
- Use dot notation to query a nested column.  
E.g. `customer.name` refers to `name` field in `customer` column.

## Repeated Fields



- **ARRAY** data type is an ordered list of zero or more elements of the same data type.  
For e.g. `product` is an **ARRAY** of **STRUCT** here.
- Use `UNNEST()` to flatten the repeated data and `OFFSET/ORDINAL` to access individual element

## Nested Repeated Fields



- Combining nested and repeated fields denormalizes a 1:many relationship without joins.
- Use dot notation to query a nested column and `UNNEST()` to flatten the repeated data.

Transaction Fact					
Order Id	timestamp	CustomerId	sku	quantity	price
1000001	12/18/2017 15:02:00	65401	ABC123456	3	36.3
1000001	12/18/2017 15:02:00	65401	TBL535522	1	878.4
1000001	12/18/2017 15:02:00	65401	CHR762222	6	435.6
1000002	12/16/2017 11:34:00	74682	GCH635354	4	345.7
1000002	12/16/2017 11:34:00	74682	GRD828822	2	9.5

Customer Dimension		
CustomerId	CustomerName	Location
65401	John Doe	Faraway
74682	Jane Michaels	Nearland
63636	Jose Carlos	Nearland


Product Dimension	
sku	description
ABC123456	Redwood 8x4
TBL535522	Sapient Table
CHR762222	Cherrywood Chair
GCH635354	Garden chairs
GRD828822	Ceramic Pots

# Summary

- BigQuery architecture is designed for petabyte-scale querying performance



Tables are broken into pieces, called shards, to allow for scalability



BigQuery uses compressed column-based storage for fast retrieval



Structs and arrays are data type containers that are foundational to repeated fields

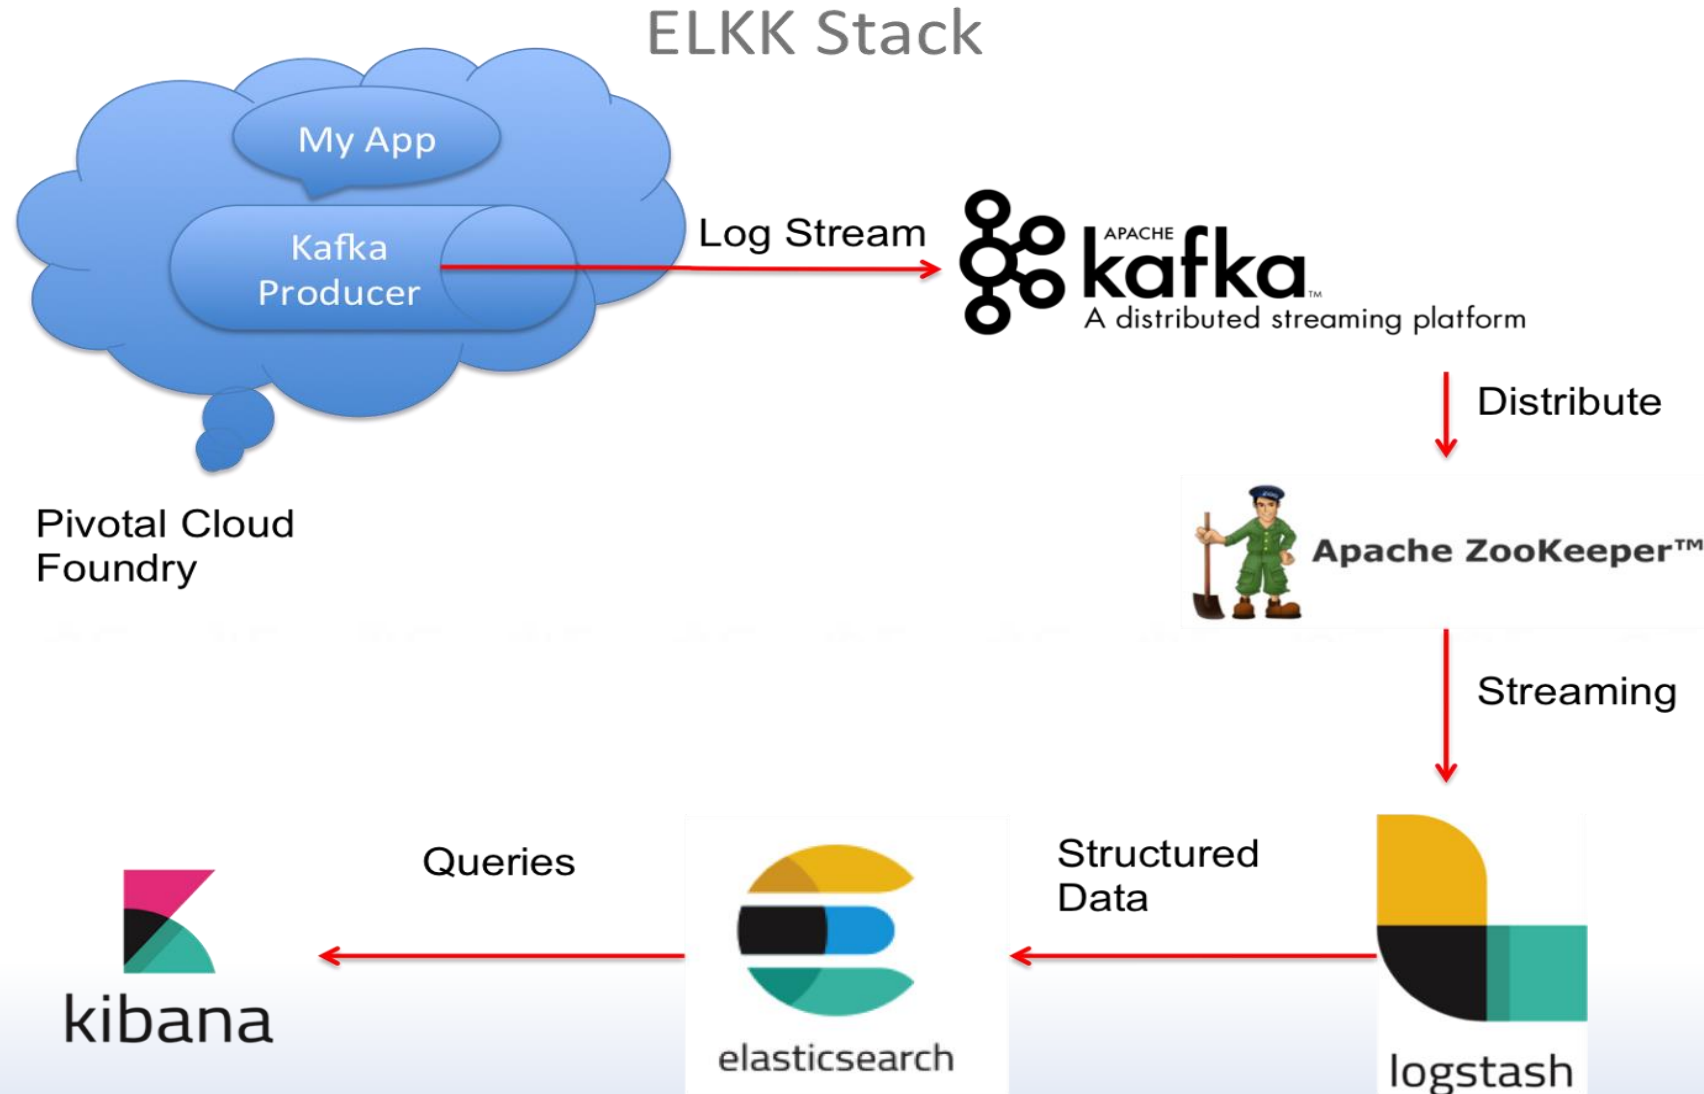
Row	date	top_articles.title
1	2010-08-23	Why GNU grep is Fast
		Readme Driven Development
2	2010-04-26	Police raid Gizmodo editor's house
		Not even in South Park?
3	2009-09-15	Learning Advanced JavaScript
		Sub-pixel re-workings of YouTube and BBC favicons

Tables with repeated fields are conceptually like pre-joined tables




# Introduction to ElasticSearch

# Real-time Data Monitoring using Kafka, Logstash, Elasticsearch and Kibana

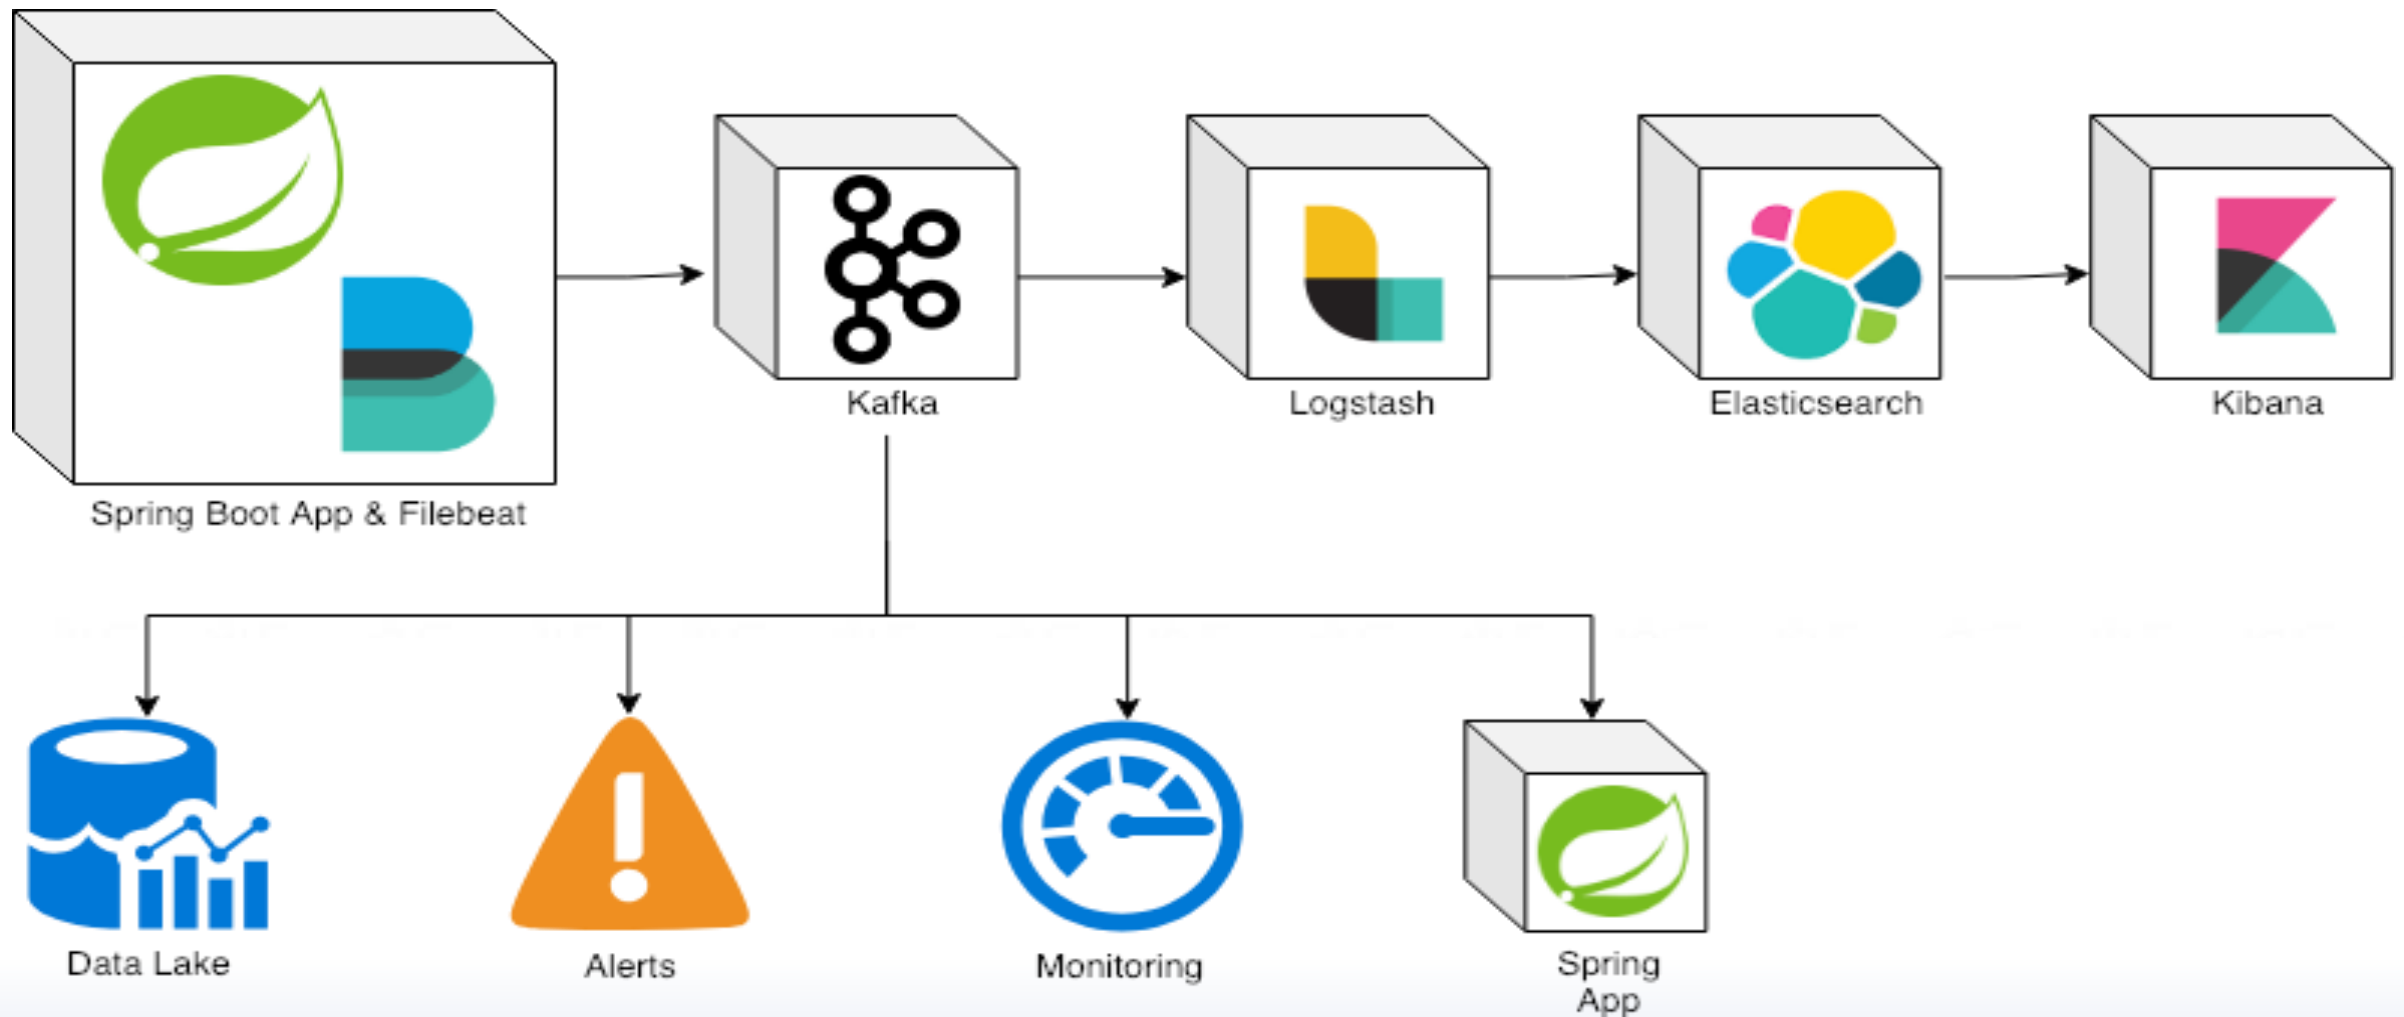


# ELK

- Elasticsearch
- Logstash
- Kibana



# Full log monitoring ecosystem



# Kibana

D

Dashboard / [Filebeat System] Syslog dashboard ECS

Full screen

Share

Clone

Edit

📅

▼

Search

KQL

📅

▼

Today

Show dates

↻

Refresh

☰

—

+ Add filter

Dashboards [Filebeat System] ECS

[Syslog](#) | [Sudo commands](#) | [SSH logs](#) | [New users and groups](#)

Syslog events by hostname [Filebeat System] ECS

Count

@timestamp per 30 minutes

Syslog hostnames and processes [Filebeat System] ECS

Rhodas-MBP
Rhodas-MacBook-P...
syslogd
com.apple.xpc.laun...
Google Chrome
com.apple.xpc.laun...
com.apple.xpc.laun...
xpcproxy
CptHost
awdd

Syslog logs [Filebeat System] ECS

1–50 of 323

<

>

Time	host.hostname	process.name	message
> May 8, 2020 @ 23:18:11.000	Rhodas-MBP	Preview	objc[47426]: Class FIFinderSyncExtensionHost is implemented in both /System/Library/PrivateFrameworks/FinderKit.framework/Versions/A/FinderKit (0x7fff981da3d8) and /System/Library/PrivateFrameworks/FileProvider.framework/OverrideBundles/FinderSyncCollaborationFileProviderOverride.bundle/Contents/MacOS/FinderSyncCollaborationFileProviderOverride (0x106f7ef50). One of the two will be used. Which one is undefined.
> May 8, 2020 @ 23:18:11.000	Rhodas-MBP	Preview	assertion failed: 18G103: libxpc.dylib + 90677 [7DEE2300-6D8E-3C00-9C63-E3E80D56B0C4]: 0x89

# What is Elasticsearch

- Document (Json) oriented search engine
- Distributed
- Horizontally scalable and highly available
- RESTful API
- Built on Lucene search engine library
- It is used for full-text search, analytics.



# ElasticSearch

- ES has become de facto *fast search* solution
- Examples
  - GitHub uses ES to query 130 billion lines of code
  - Wikipedia uses ES to provide full-text search with highlighted search snippets
  - StackOverflow combines full-text search with geolocation queries and uses more-like-this to find related questions and answers

# History



Shay Benon @kimchy

**Elasticsearch** released in February 2010.  
Worked on this for 6 years (started with compass)  
Now part of <http://elastic.co> commercial offerings



Doug Cutting @cutting

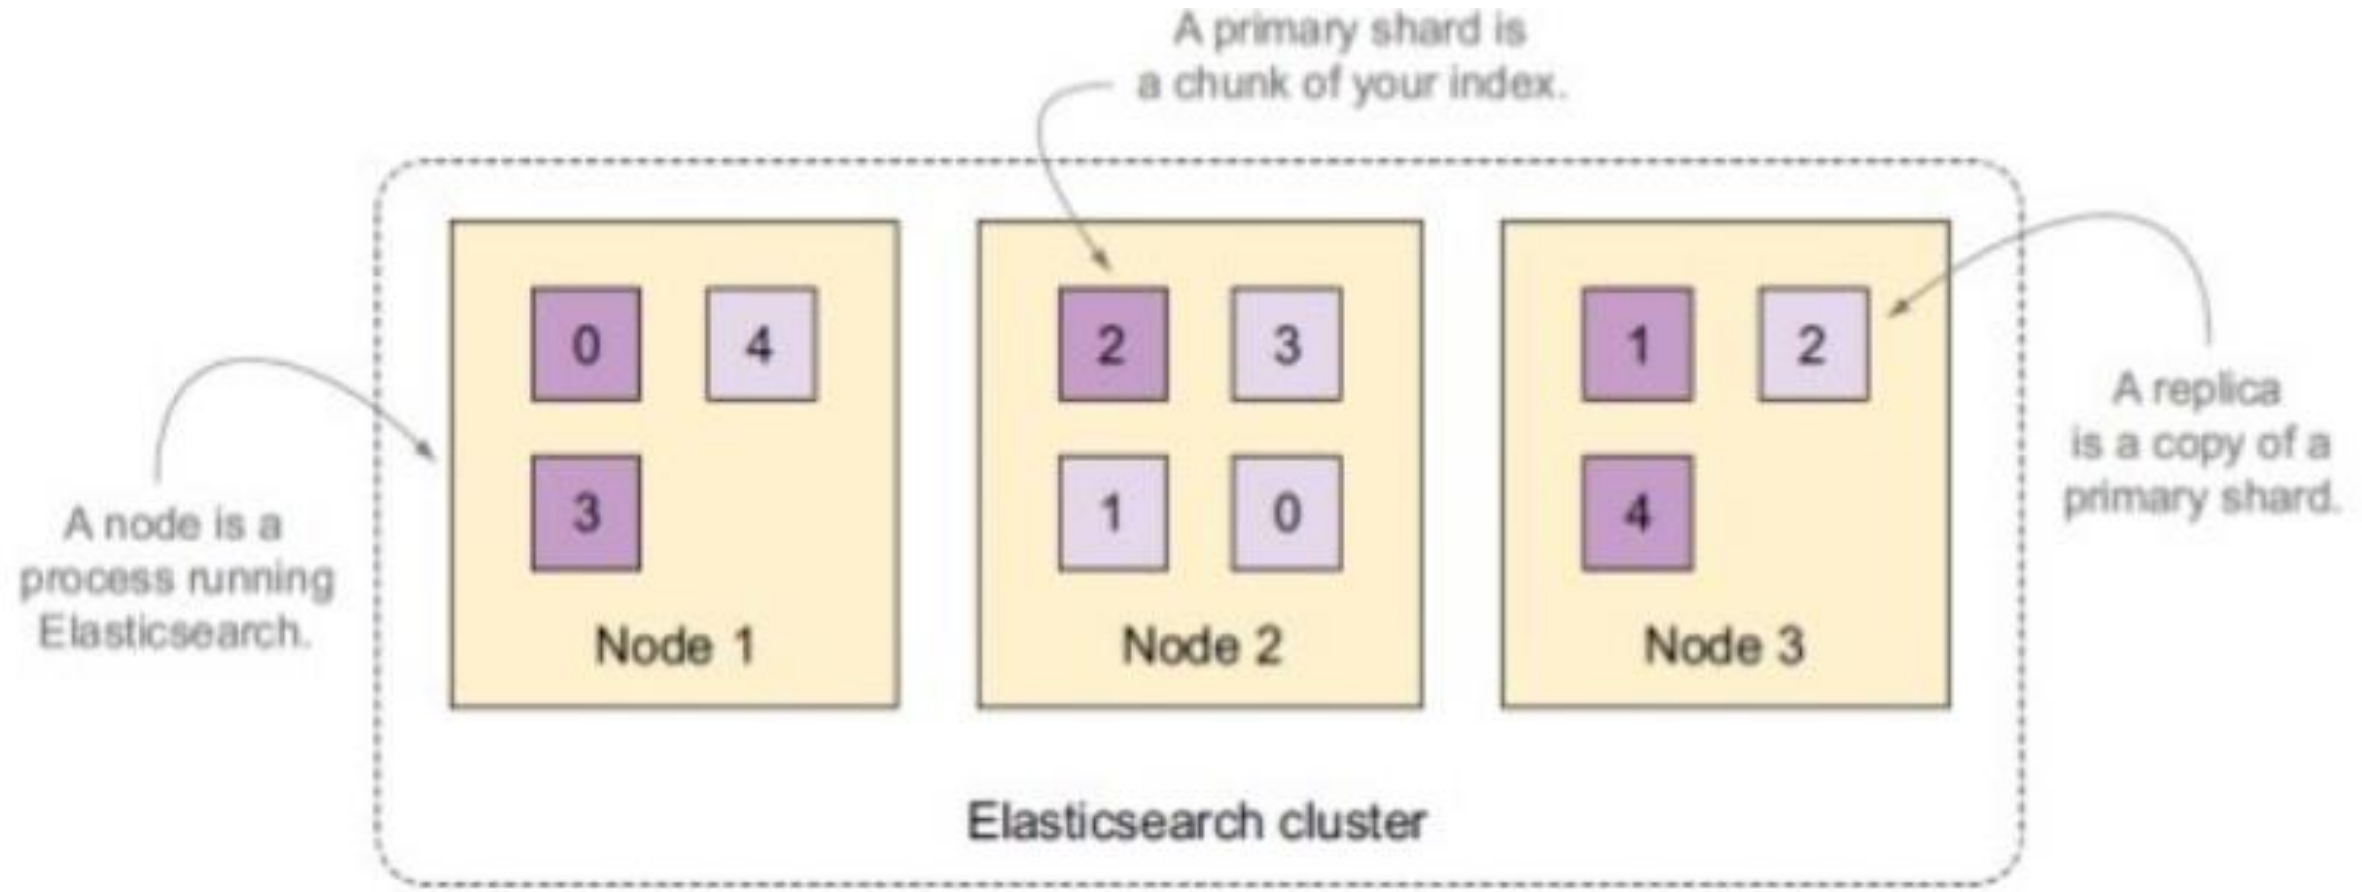
Started **Lucene** in 1999, released under apache in 2005.  
Now part of cloudera supporting rival solution solr and commercial offerings

# Building Blocks

Term	Description ( ~analogy with relational database)
Cluster	~Database cluster Group of nodes
Node	~Instance of database  A JVM process, usually a machine
Index	~Database schema Hosts mapping types and their definitions contains many shards
Mapping Type	~Database Table Field description, indexing requirements
Document	~Database row Json document.
Shard	A Lucene index. Scalable unit and heart of search engine ( <i>primary and replica</i> )

Relational DB	Elasticsearch 6
Table	Index
Row	Document
Column	Field

# Physical Layout



# Inverted Index

•**Step 1** - we need tokenize those docs into terms with a tokenizer. So, let's say, we use Tokenizer A and get following results:

Doc1: Harry, Potter, And, The, Half,

Doc2: Harry, Potter, And, The, Deathly

•**Step 2** - Build inverted index

Harry -> Doc1, Doc2

Potter -> Doc1, Doc2

And -> Doc1, Doc2

The -> Doc1, Doc2

Half -> Doc1

Blood -> Doc1

Prince -> Doc1

Deathly -> Doc2

Hallows -> Doc2

## To query/search

•**Step 1** - We also need tokenize search words at first. Such as, our search words are Harry Potter.

And you have two tokenizer to choose.

Tokenizer A is same with the one we use during indexing, will tokenize our words into two terms: Harry and Potter.

•**Step 2** - Do query

If you choose Tokenizer A, you get Harry and Potter, both of them are in our inverted index, then you can get search results: Doc1 and Doc2.

# Lucene Inverted Index

A shard is a Lucene index.  
get-together0 shard

Inverted index		
Term	Document	Frequency
elasticsearch	id1	1 occurrence: id1->1 time
denver	id1, id3	3 occurrences: id1->1 time, id3->2 times
clojure	id2, id3	5 occurrences: id2->2 times, id3->3 times
data	id2	2 occurrences: id2->2 times

## ES: Lucene +

- Distributed
- Transaction Log
  - The translog **stores all operations that are not yet safely persisted in Lucene**. Although these operations are available for reads, they will need to be replayed if the shard was stopped and had to be recovered.

Simplifies shared relocation/recovery

- Query DSL
  - Provides set of grammar for search syntax



# Index>Type>Document>Field

Vehicles(index)

Car(document 1)

Car (document 2)

Car (document n)

```
{  
  "type": "car",  
  "documents": [  
    {  
      "id": 23134,  
      "make": "Honda",  
      "color": "red",  
      "mileage": 8000  
    },  
    {  
      "id": 12334,  
      "make": "Ford",  
      "color": "white",  
      "mileage": 5000  
    }  
  ]  
}
```

# index docs

Inserting = indexing

```
PUT /{index}/{type}/{id}
{
  "field1": "value1",
  "field2": "value2",
  ...
}
```

```
PUT /vehicles/car/123
{
  "make": "Ford",
  "mileage": 5000,
  "color": "red"
}
```

# Search Types

## Count

Returns no hits, only total count matching the query

## Scan

Allows to iterate over large amounts of data using a cursor to paginate and memory efficient

## Search

General search

# Query DSL (Domain Specific Language)

Get /courses/\_search

```
{
  "query": {
    "match_all": {}
  }
}
```

Get /courses/\_search

```
{
  "query": {
    "match": {
      "name": "computer"
    }
  }
}
```

```
{
  "took": 0,
  "timed_out": false,
  "_shards": {
    "total": 1,
    "successful": 1,
    "failed": 0
  },
  "hits": {
    "total": 1,
    "max_score": 0.81233,
    "hits": [
      {
        "_index": "courses",
        "_type": "classroom",
        "_id": "3",
        "_score": 0.8243,
        "_source": {
          "name": "Computer Eng 101",
          "room": "ENG202",
          "student enrolled": 23
        }
      }
    ]
  }
}
```

# Match query

The match query is the standard query for performing a full-text search.

```
GET /_search
{
  "query": {
    "match": {
      "message": {
        "query": "this is a test"
      }
    }
  }
}
```

## Getting a document

Request:

GET test/cities/1?pretty

Response:

```
{
  "_index": "test",
  "_type": "cities",
  "_id": "1",
  "_version": 1,
  "found": true,
  "_source": {
    "rank": 3,
    "city": "Hyderabad",
    "state": "Telangana",
    "population2014": 7750000,
    "land_area": 625,
    "location": {
      "lat": 17.37,
      "lon": 78.48
    },
    "abbreviation": "Hyd"
  }
}
```

# Document Metadata Fields

- ***\_id*** - The id of the document
- ***\_type*** - The document type
- ***\_source*** - enabled Stores the original document that was indexed
- ***\_all*** enabled Indexes all values of all document fields
- ***\_timestamp*** disabled timestamp associated with the document
- ***\_ttl*** disabled optionally defines an expiration time
- ***\_size*** disabled indexes the size of the uncompressed

# Query DSL (Domain Specific Language)

```
% curl 'localhost:9200/get-together/_search?sort=date:asc&_source=title,date'
{
  "index": "get-together",
  "type": "event",
  "id": "114",
  "score": null,
  "source": {
    "date": "2013-09-09T18:30",
    "title": "Using Hadoop with Elasticsearch"
  },
  "sort": [
    1378751400000
  ]
},
```

← Show one hit of the response.

Request matching all documents but return the default first 10 of all results ordered by date in ascending order. You want only two fields in the response: title and date.

← The filtered `_source` document now contains only filtered fields.

← The score is null; you're using a sort and therefore no score is calculated.

<https://www.elastic.co/guide/en/elasticsearch/reference/current/query-dsl.html>



# Aggregation

- An aggregation summarizes your data as metrics, statistics, or other analytics. Aggregations help you answer questions like:
- What's the average load time for my website?

```
GET /my-index-000001/_search
{
  "aggs": {
    "my-agg-name": {
      "terms": {
        "field": "my-field"
      }
    }
  }
}
```

# demo

Web application with backend DB as ES, autocomplete.

<https://www.youtube.com/watch?v=hVSC4ZNiVdA>

Identifying anomalies and forecasting with ES

<https://www.youtube.com/watch?v=wJVgh5knV4E>

the ELK Stack: Elasticsearch, Kibana, Beats, and Logstash

<https://www.youtube.com/watch?v=DRQJHOOstY0>

Application Performance Monitoring (APM) with Elasticsearch, Elastic Stack

<https://www.youtube.com/watch?v=2sdOvuLiBb8>

Any volunteer to build a **full stack ML** with ELK, node.js and React.js?