Introduction to

Slide Source: 

https://go.neo4j.com/rs/710-RRC-335/images/Graph_Databases_for_Beginners.pdf

Neo4j is the creator of

a highly scalable, ***native*** graph database.

Neo4j gives any organization the ability to ***leverage connections in data — in real-time*** *to create value*

Our core belief is — *connections between data are as important as the data itself*

Use of data connections has created industry leaders

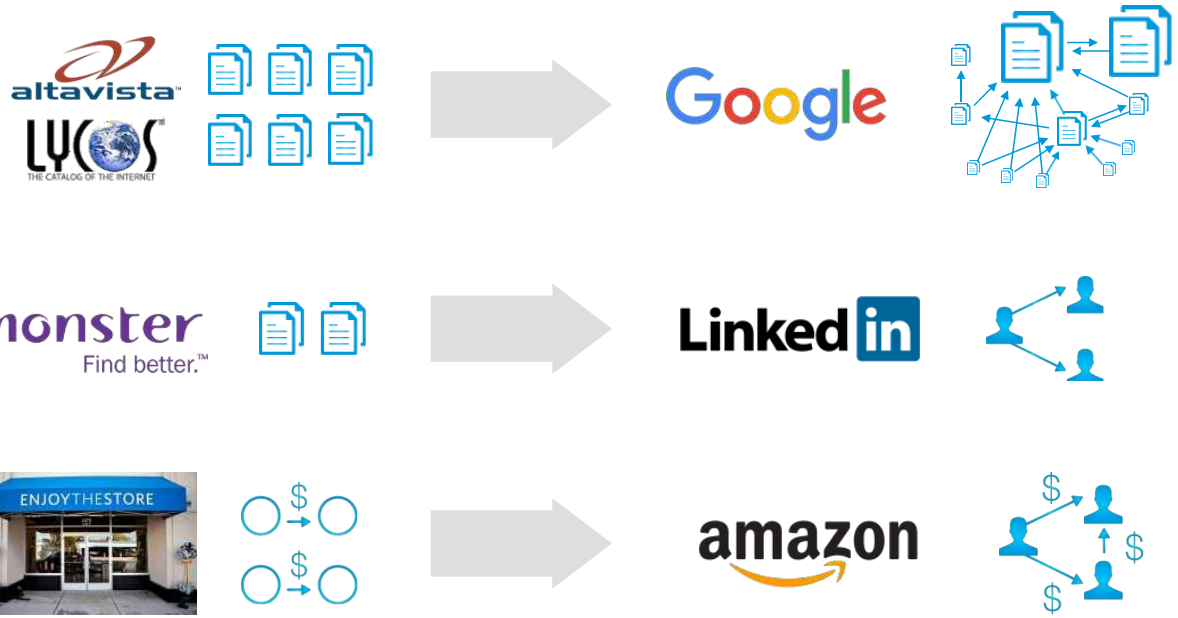

Linked Connections

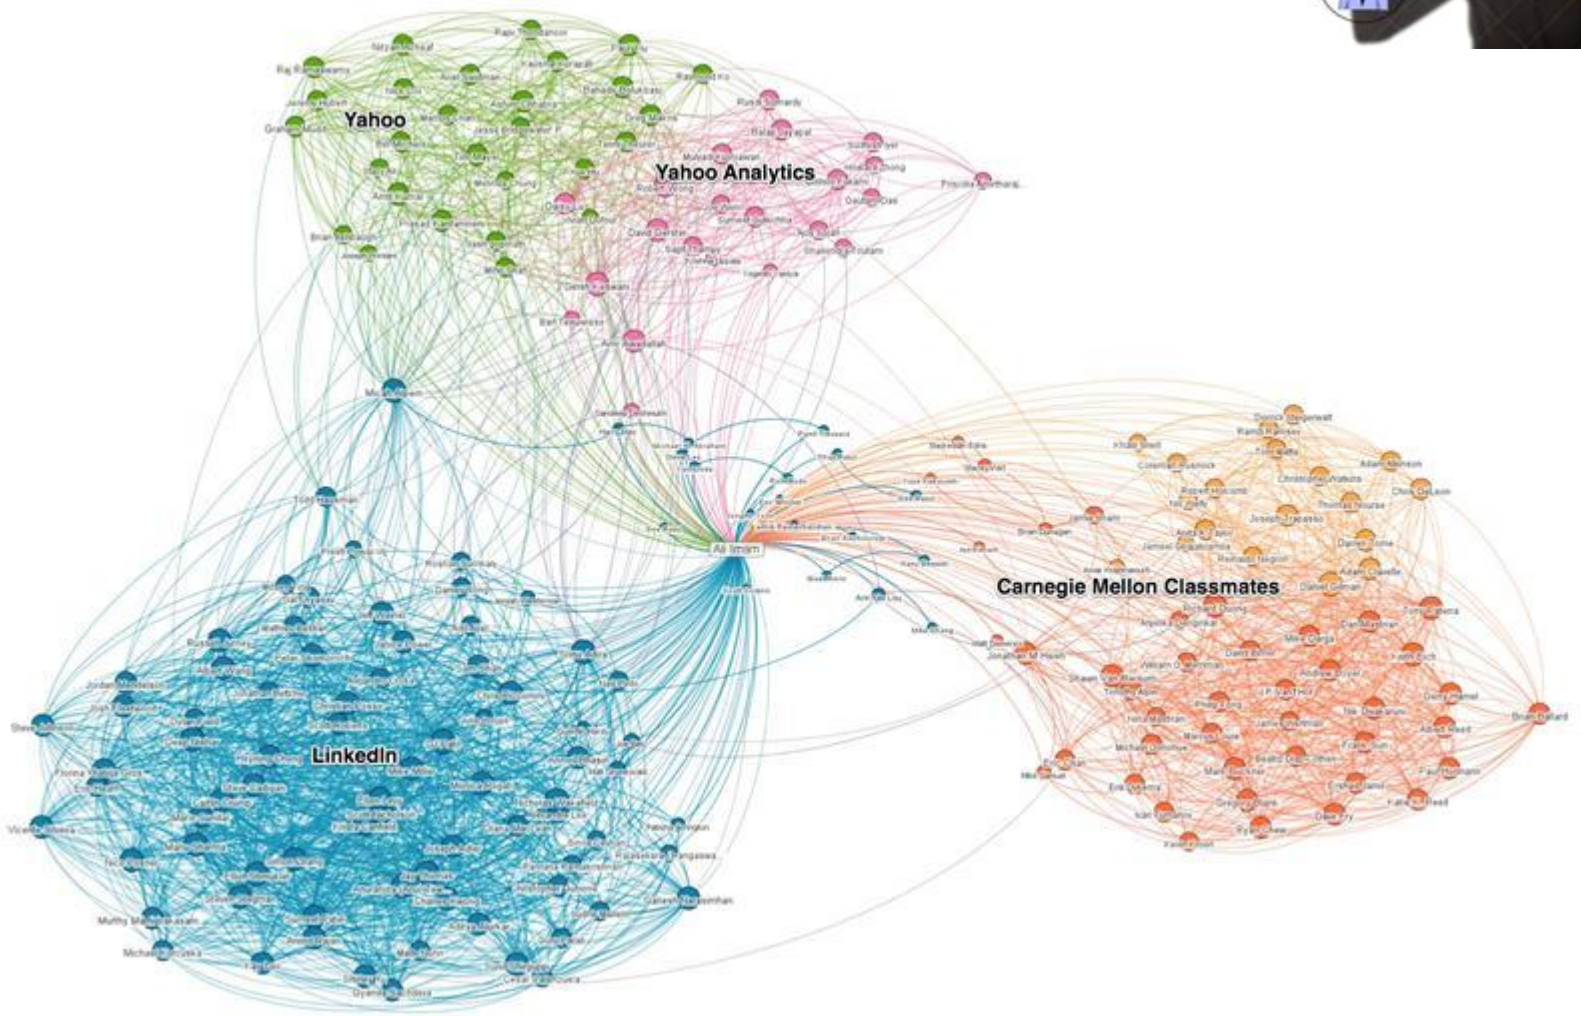

Today, as processes get digitized and interconnected - we see the emergence of a *"connected enterprise"*

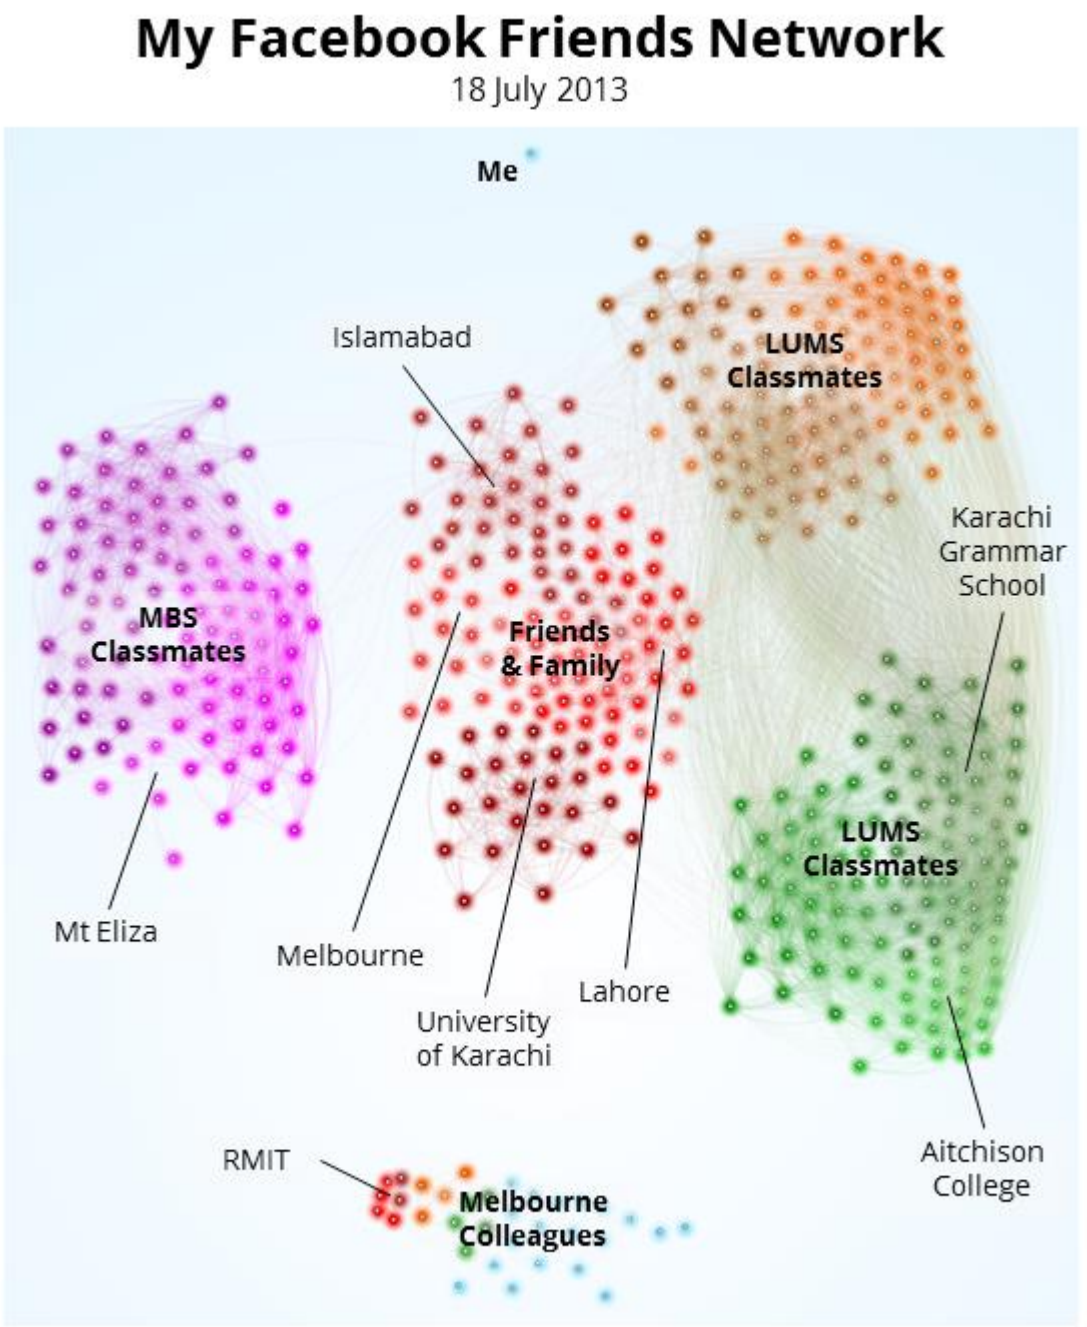

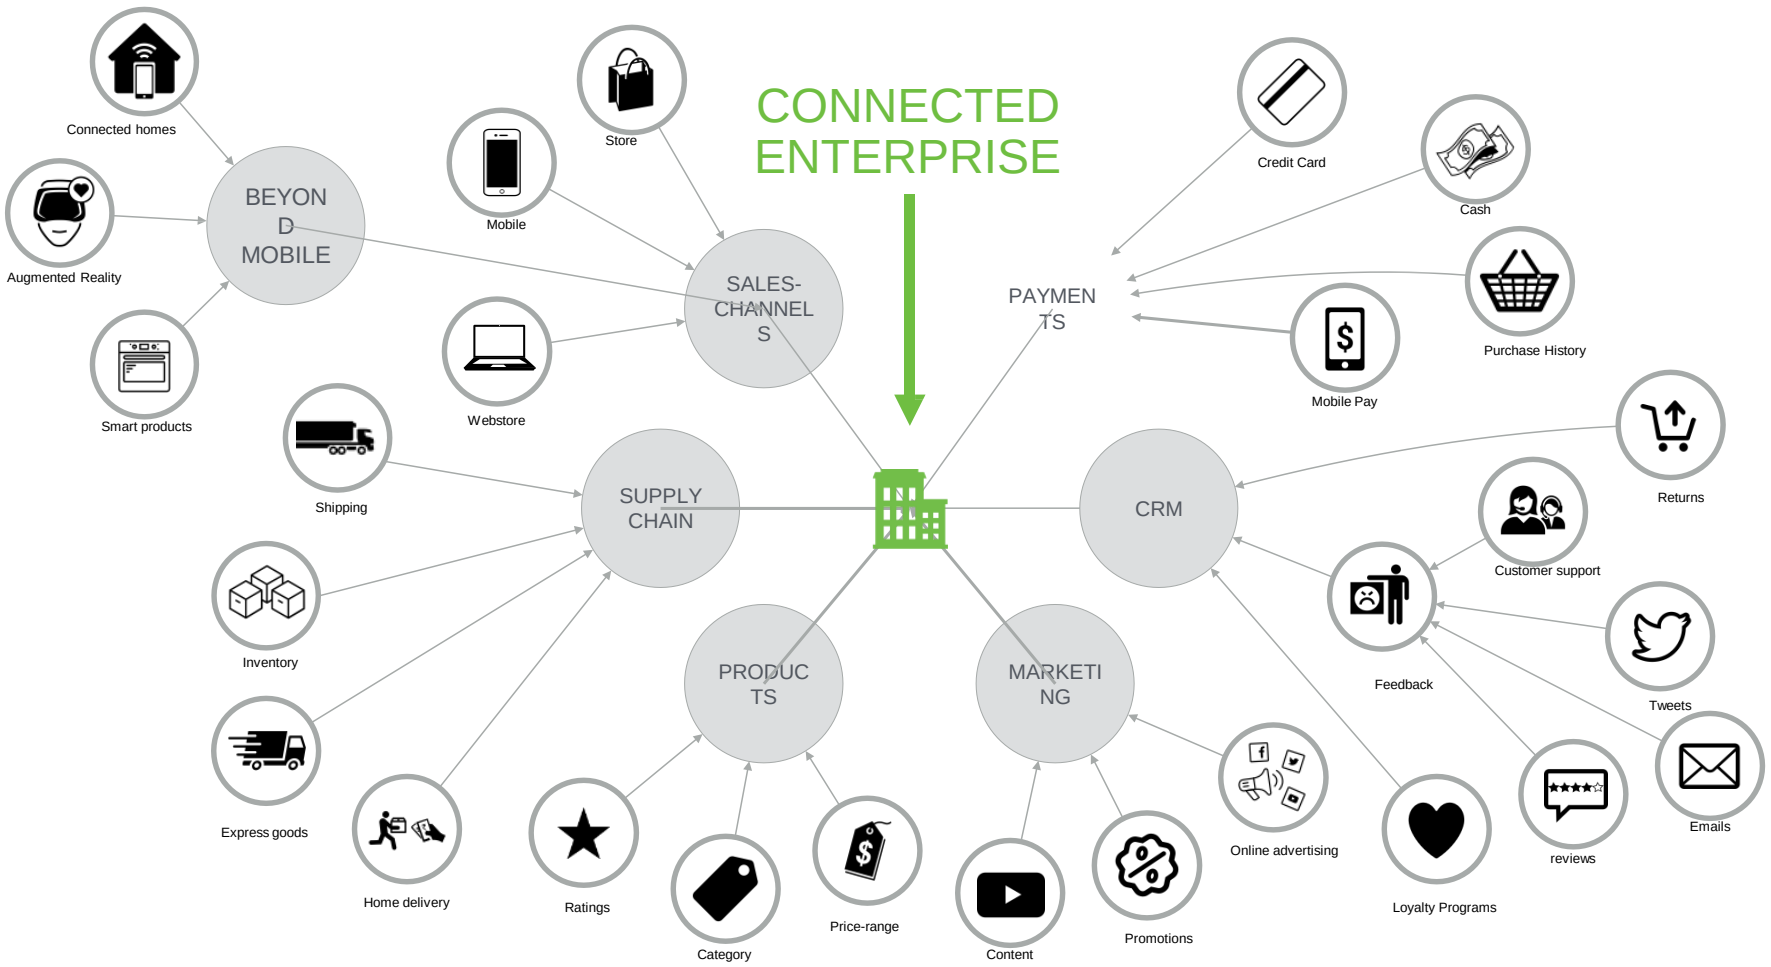
CONNECTED ENTERPRISE
BEYOND MOBILE
Connected homes
Augmented Reality
Smart products
SALES-CHANNELS
Mobile
Store
Webstore
PAYMENTS
Credit Card
Cash
Purchase History
Mobile Pay
SUPPLY CHAIN
Shipping
Inventory
Express goods
Home delivery
CRM
Returns
Customer support
Feedback
Tweets
Emails
reviews
Loyalty Programs
PRODUCTS
Ratings
Category
Price-range
MARKETING
Online advertising
Promotions
Content

Relationship Queries Strain Traditional Databases

A single query can touch a ***lot of data***

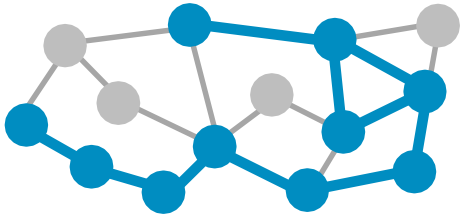

Queries can take non-sequential, ***arbitrary paths*** through data

Real-time queries need speed and ***consistent response times***

Queries must ***run reliably*** with ***consistent results***

What you (probably) already know:

SQL Join Hell (1)

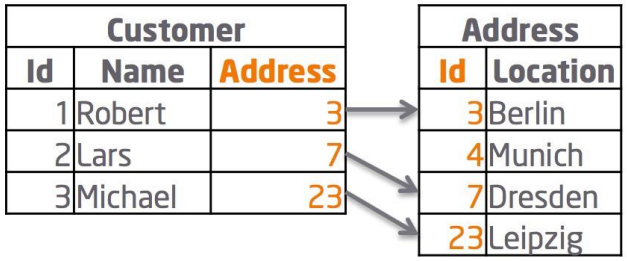

Customer		
Id	Name	Address
1	Robert	3
2	Lars	7
3	Michael	23

Address	
Id	Location
3	Berlin
4	Munich
7	Dresden
23	Leipzig

1:1 Relationship

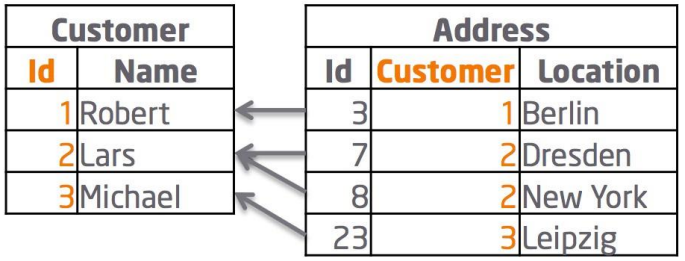

Customer	
Id	Name
1	Robert
2	Lars
3	Michael

Address		
Id	Customer	Location
3	1	Berlin
7	2	Dresden
8	2	New York
23	3	Leipzig

1:n Relationship

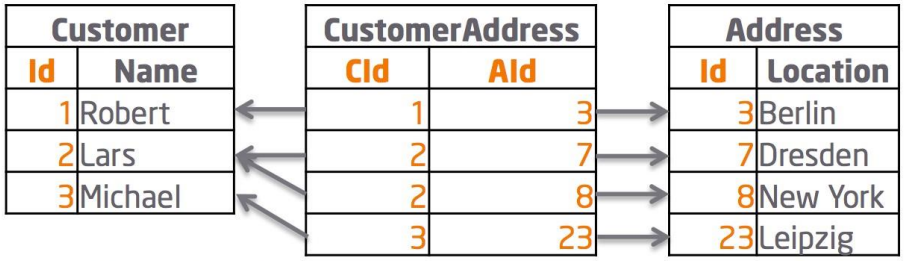

Customer	
Id	Name
1	Robert
2	Lars
3	Michael

CustomerAddress	
CId	AId
1	3
2	7
2	8
3	23

Address	
Id	Location
3	Berlin
7	Dresden
8	New York
23	Leipzig

m:n Relationship

The Problem

1. Joins are executed **every time** you query the relationship

2. Executing a Join means to **search** for a key

3. **B-Tree Index: O(log(n))**
   Your data grows by 10x, your time goes up by one step on each Join

4. More Data = More Searches
   Slower Performance

Relational Databases can't handle Relationships

1 **Wrong Model**
They cannot model or store relationships without complexity

2 **Degraded Performance**
Speed plummets as data grows and as the number of joins grows

3 **Wrong Language**
SQL was built with Set Theory in mind, not Graph Theory

4 **Not Flexible**
New types of data and relationships require schema redesign

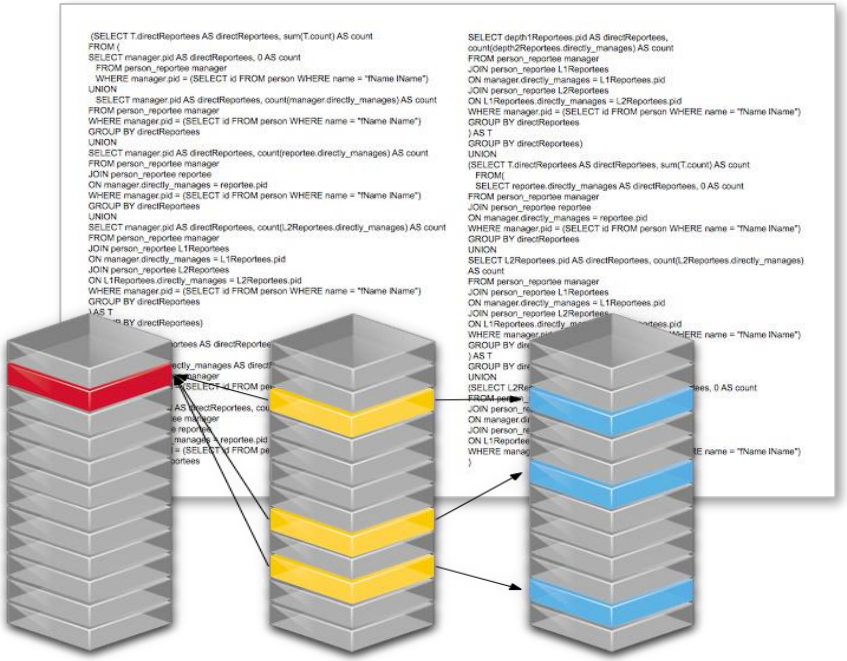

Relational Versus Graph Models

Relational Model

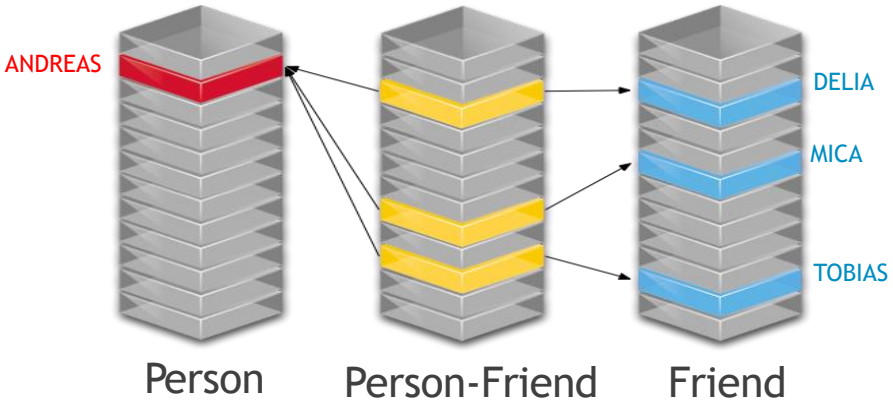

Graph Model

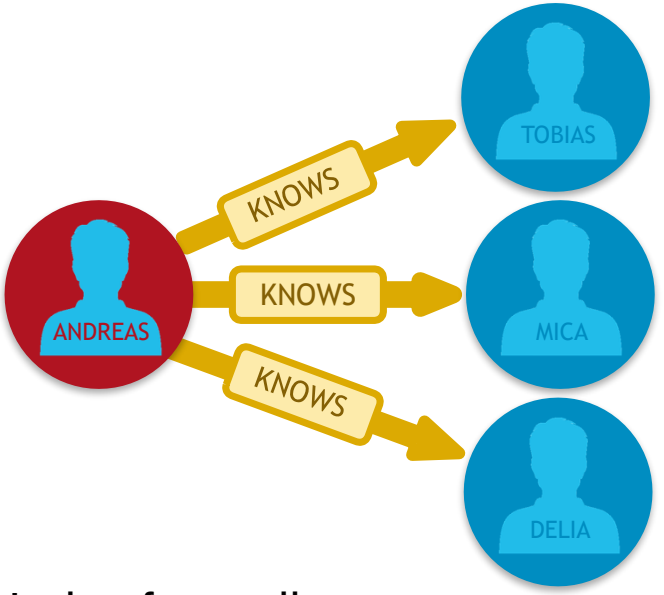

Index free adjacency

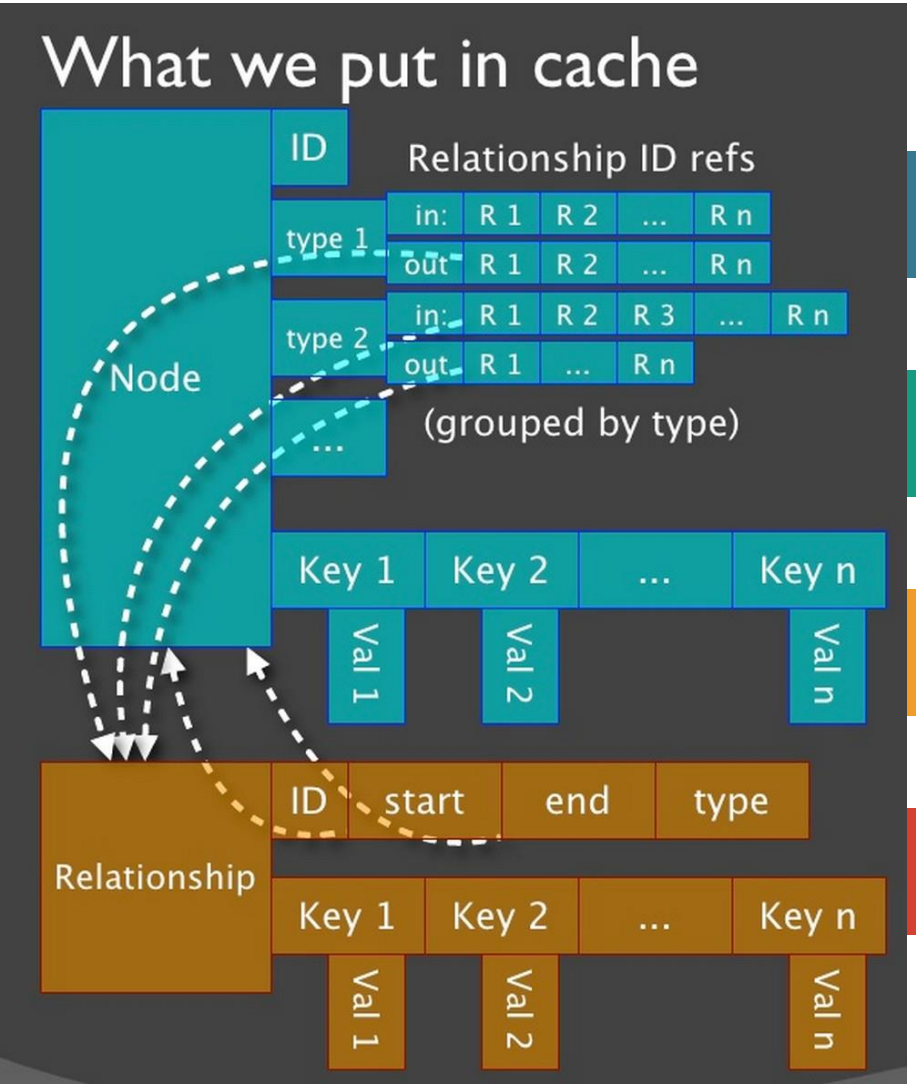

1 Pointers instead of Lookups

2 Fixed Sized Records

3 "Joins" on Creation

4 Spin through this data structure

Real Time Query Performance

Remains steady as database grows

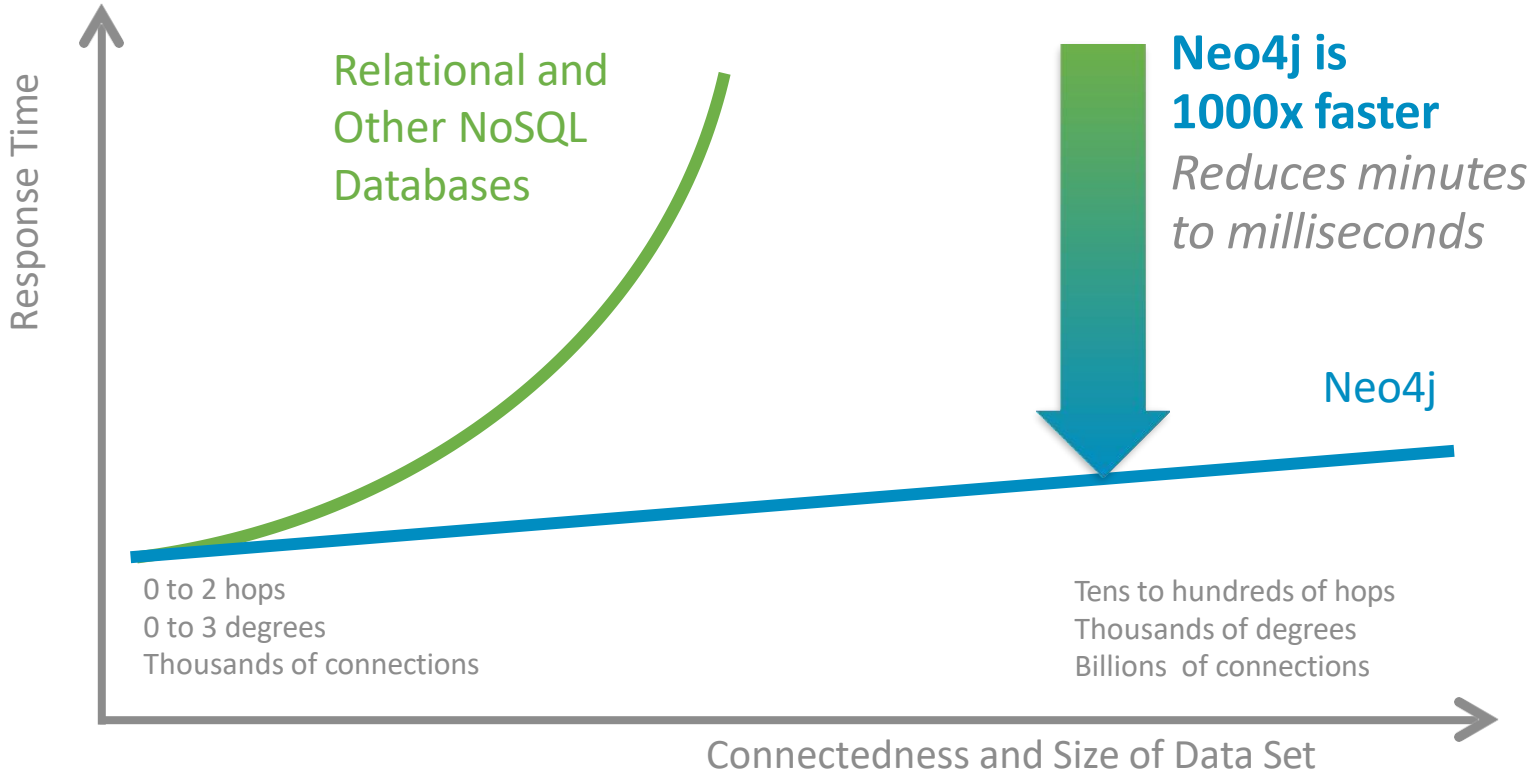

Reimagine your Data as a Graph

1 **Right Model**
Graphs simplify how you think

2 **Better Performance**
Query relationships in real time

3 **Right Language**
Cypher was purpose built for Graphs

4 **Flexible and Consistent**
Evolve your schema seamlessly while keeping transactions

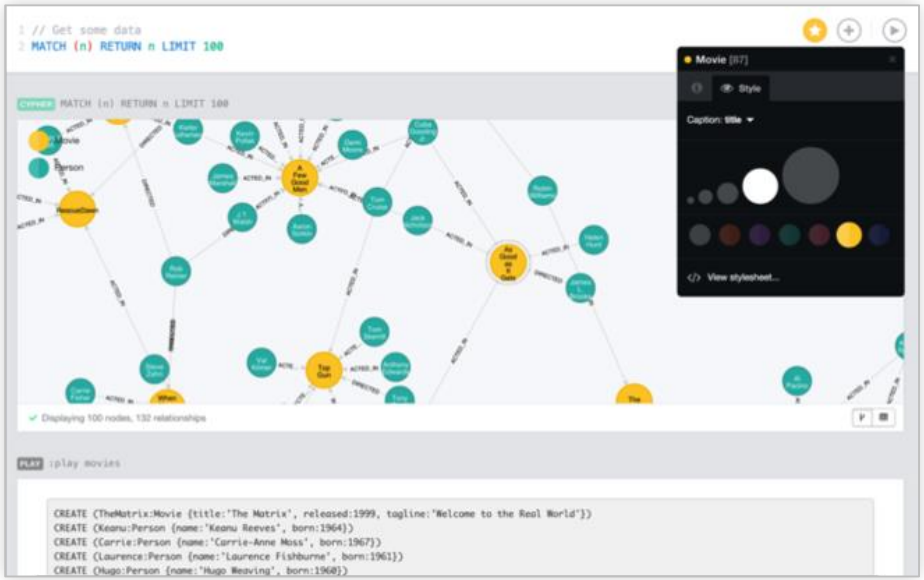

Agile, High Performance and Scalable without Sacrifice

Neo4j - The #1 Database for Connected Data

Neo4j is an ***enterprise-grade native graph database*** that enables you to:

- **Store and query** data relationships
- **Traverse** any levels of depth on real-time
- **Add and connect** new data on the fly

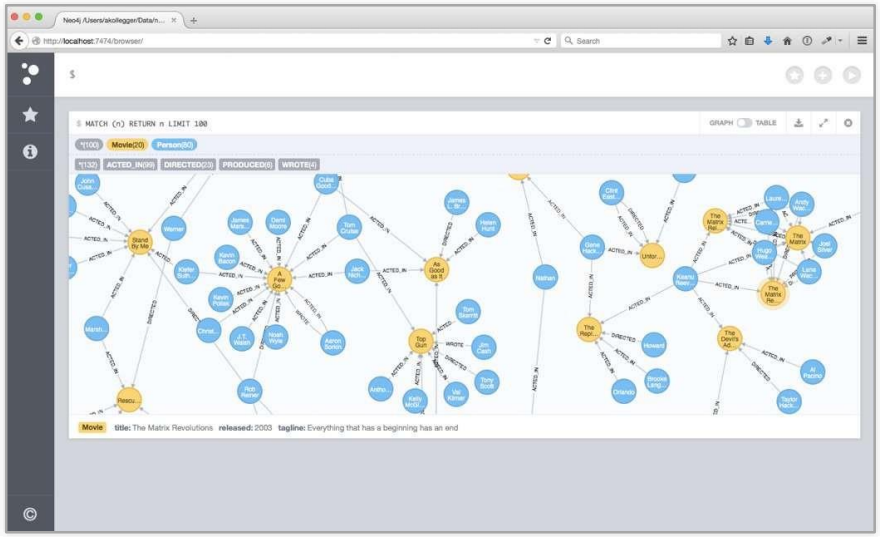

Designed, built and tested natively for graphs from the start to ensure:

- Performance
- ACID Transactions
- Agility
- Developer Productivity
- Hardware Efficiency

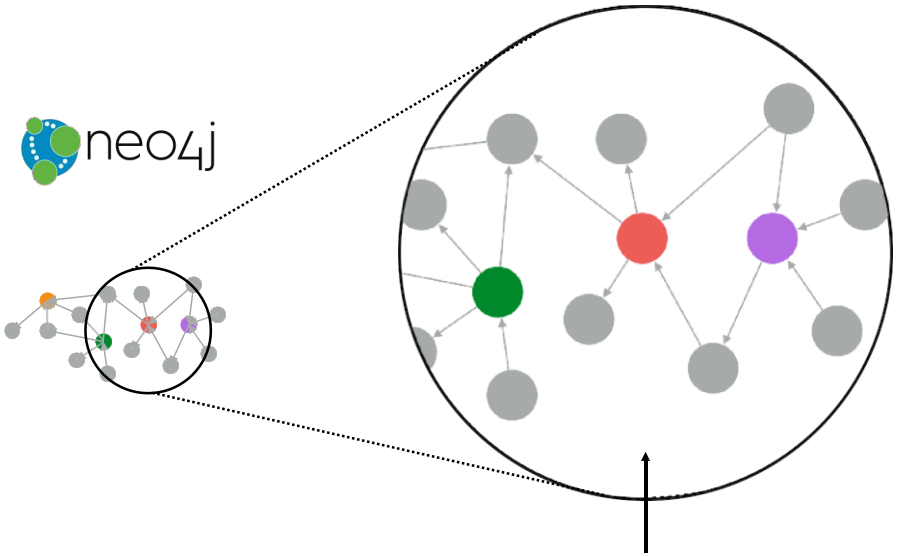

Index-free adjacency ensures lightning-fast retrieval of data and relationships

Index free adjacency:

Unlike other database models Neo4j ***connects*** data as it stores it

Neo4j Advantage - Performance

Neo4j: Native Graph from the Start

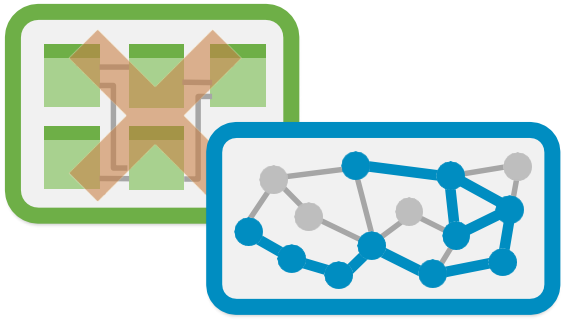

Native graph storage

Optimized for real-time reads and ACID writes

- **Relationships stored as physical objects, eliminating need for joins and join tables**
- Nodes connected at write time, enabling scale-independent response times

Native graph querying

Memory structures and algorithms optimized for graphs

- Index-free adjacency enables 1M+ hops per second via in-memory pointer chasing
- Off-heap page cache improves operational robustness and scaling compared with JVM-based caches
- “Minutes to milliseconds” performance improvement

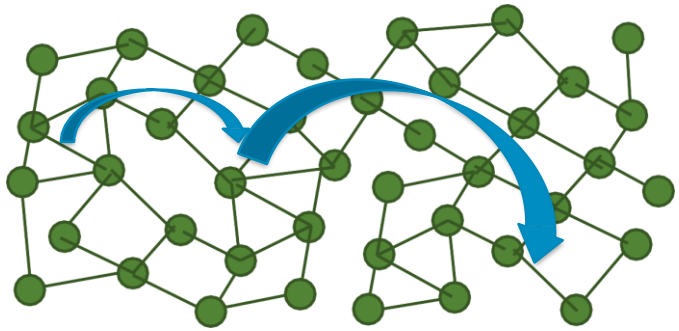

Neo4j Advantage - Performance

As a Graph

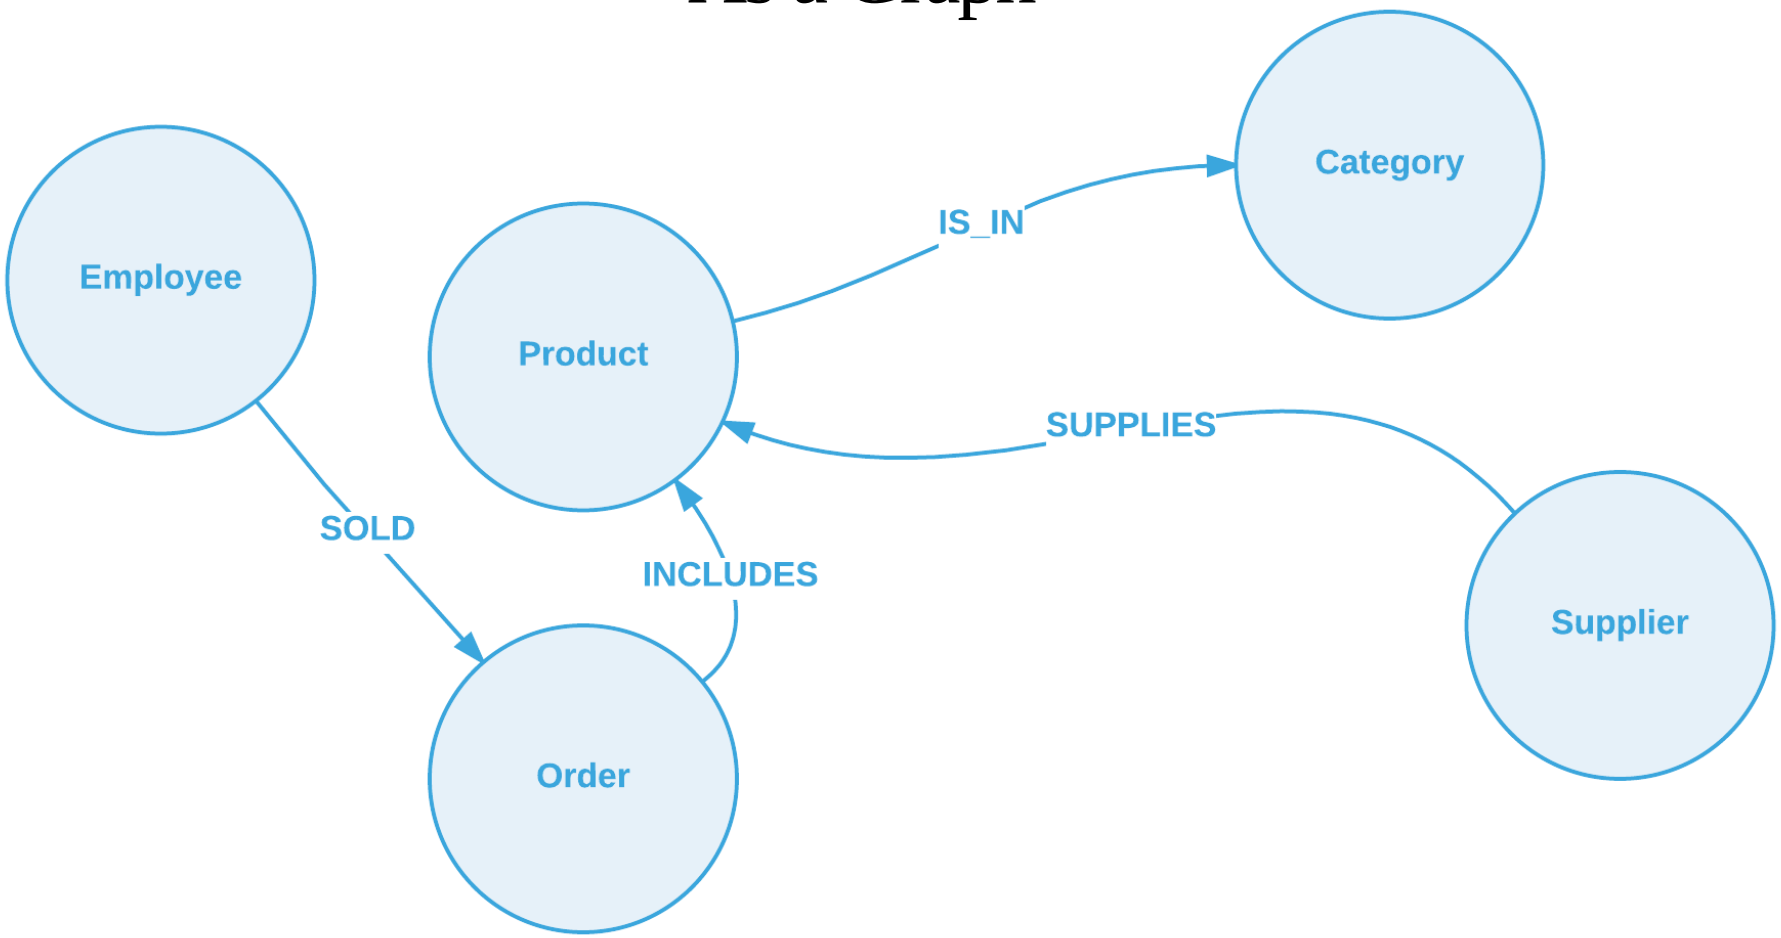
Employee
Product
Category
IS_IN
SOLD
INCLUDES
SUPPLIES
Supplier
Order

Express Complex Relationship Queries Easily with Cypher

Find all direct reports and how many people they manage, up to three levels down

Cypher Query

```
MATCH (boss)-[:MANAGES*0..3]->(sub),
      (sub)-[:MANAGES*1..3]->(report)
WHERE boss.name = “John Doe”
RETURN sub.name AS Subordinate,
  count(report) AS Total
```

SQL Query

```
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
SELECT manager.pid AS directReportees, 0 AS count
  FROM person_reportee manager
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
  SELECT manager.pid AS directReportees, count(manager.directly_manages) AS count
FROM person_reportee manager
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
UNION
SELECT manager.pid AS directReportees, count(reportee.directly_manages) AS count
FROM person_reportee manager
JOIN person_reportee reportee
ON manager.directly_manages = reportee.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
UNION
SELECT manager.pid AS directReportees, count(L2Reportees.directly_manages) AS count
FROM person_reportee manager
JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
SELECT manager.directly_manages AS directReportees, 0 AS count
FROM person_reportee manager
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
SELECT reportee.pid AS directReportees, count(reportee.directly_manages) AS count
FROM person_reportee manager
JOIN person_reportee reportee
ON manager.directly_manages = reportee.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
UNION
SELECT reportee.pid AS directReportees, count(L2Reportees.directly_manages) AS count
FROM person_reportee manager
JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
UNION
```

```
SELECT depth1Reportees.pid AS directReportees,
count(depth2Reportees.directly_manages) AS count
FROM person_reportee manager
JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
  FROM(
  SELECT reportee.directly_manages AS directReportees, 0 AS count
FROM person_reportee manager
JOIN person_reportee reportee
ON manager.directly_manages = reportee.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
UNION
SELECT L2Reportees.pid AS directReportees, count(L2Reportees.directly_manages)
AS count
FROM person_reportee manager
JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT L2Reportees.directly_manages AS directReportees, 0 AS count
FROM person_reportee manager
JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
)
```

21

Neo4j Advantage - Developer Productivity

Graph_Databases_for_Beginners.pdf (neo4j.com)

Property Graph Model

CREATE (:Employee{ firstName:“Steven”}) -[:REPORTS_TO]-> (:Employee{ firstName:“Andrew”})

LABEL PROPERTY LABEL PROPERTY

Who do people report to?

```
MATCH
  (e:Employee)<-[:REPORTS_TO]-(sub:Employee)
RETURN
   *
```

Who do people report to?

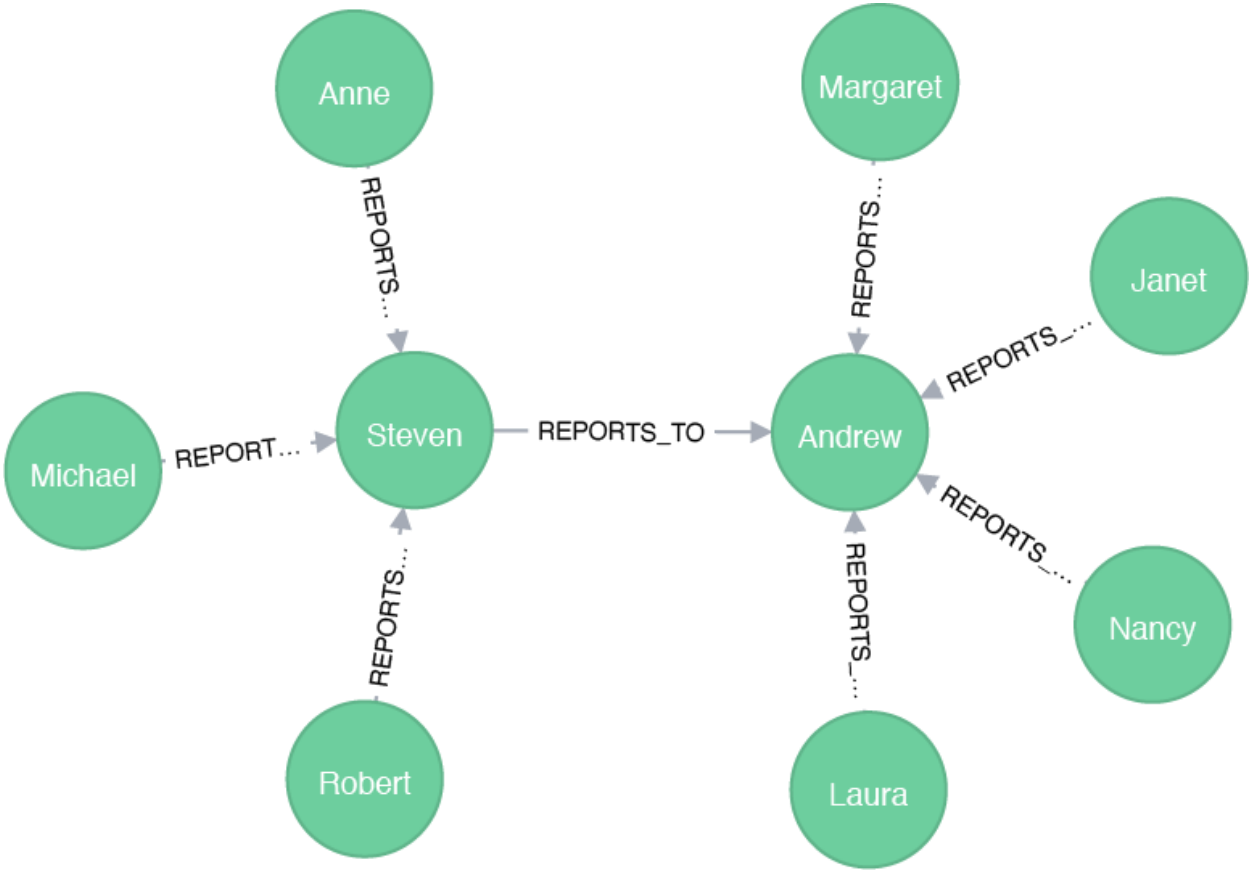

Who do people report to?

```
$ MATCH (e:Employee)<-[:REPORTS_TO]-(sub:Employee) RETURN e.employeeID AS...
```

Rows

Code

managerID	managerName	employeeID	employeeName
2	Andrew	1	Nancy
2	Andrew	3	Janet
2	Andrew	4	Margaret
2	Andrew	5	Steven
2	Andrew	8	Laura
5	Steven	9	Anne
5	Steven	6	Michael
5	Steven	7	Robert

Returned 8 rows in 111 ms.

Who does Robert report to?

```
MATCH
 p=(e:Employee)<-[:REPORTS_TO]-(sub:Employee)
WHERE
   sub.firstName = 'Robert'
RETURN
   p
```

Who do people report to?

```
MATCH
  (e:Employee)<-[:REPORTS_TO]-(sub:Employee)
RETURN
   e.employeeID AS managerID,
   e.firstName AS managerName,
   sub.employeeID AS employeeID,
   sub.firstName AS employeeName;
```

Representing Bi-Directionality

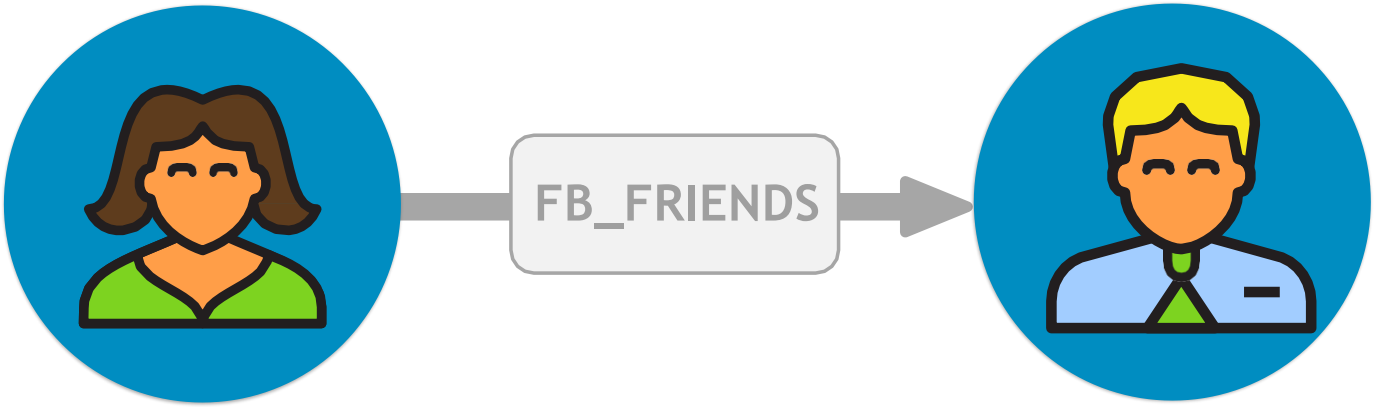

MATCH (:Person { name:“Ann”}) **- [:FB_FRIENDS] - >** (:Person { name:“Dan”})

MATCH (:Person { name:“Ann”}) **<- [:FB_FRIENDS] -** (:Person { name:“Dan”})

Modeling a Twitter Feed

Bigger Model

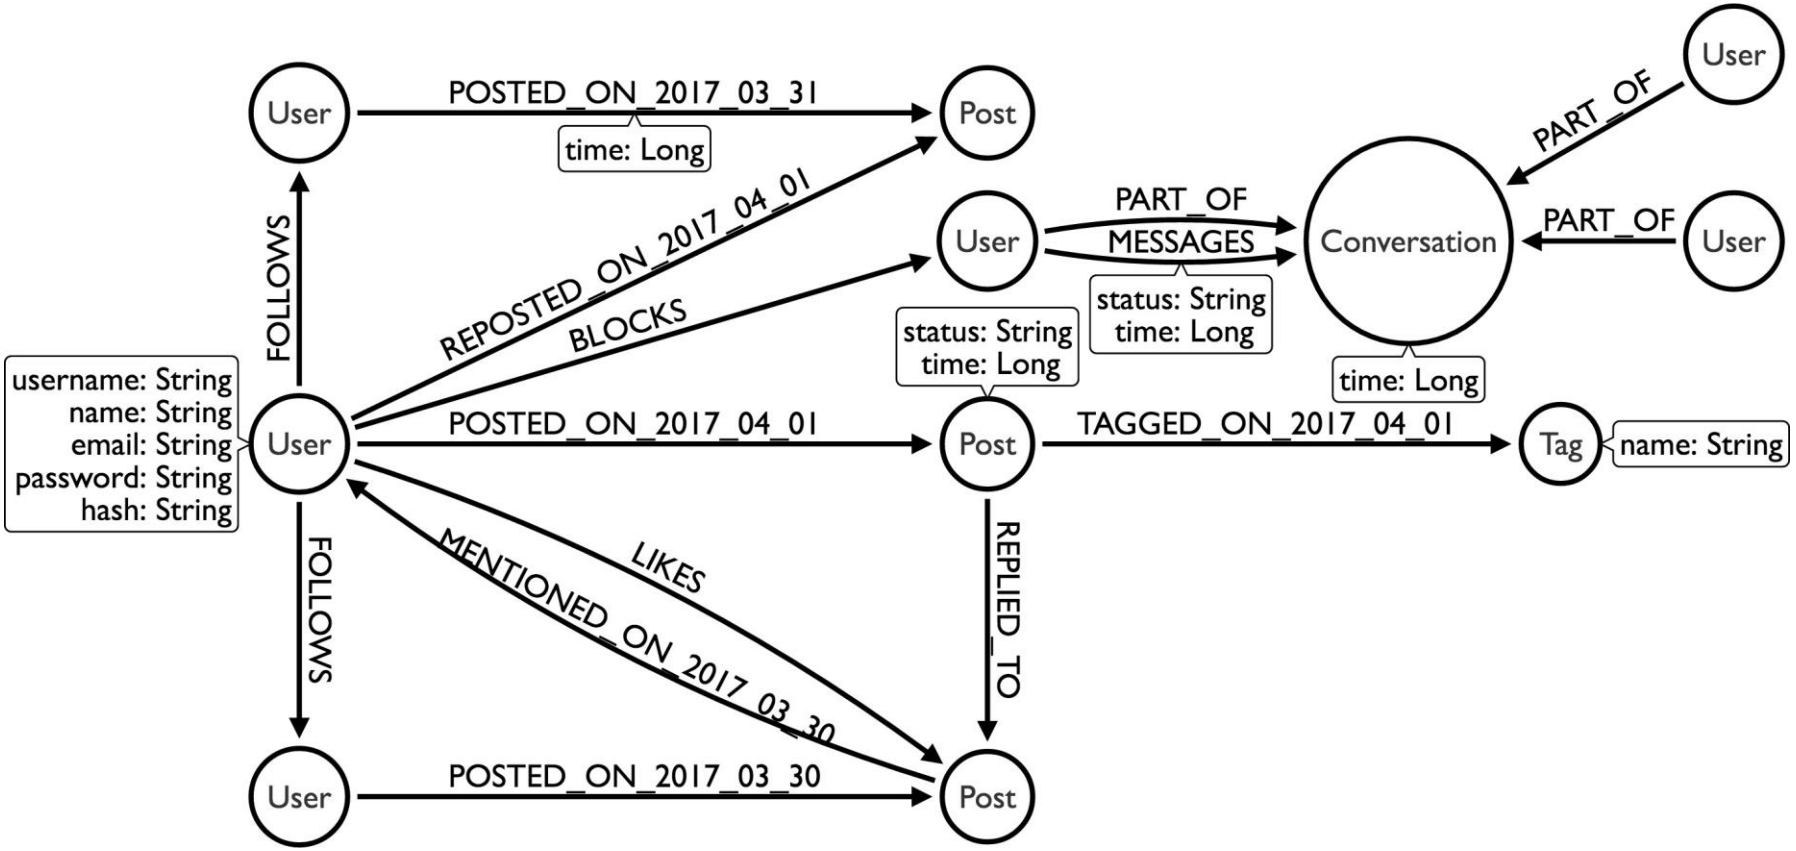

Movie Data Model

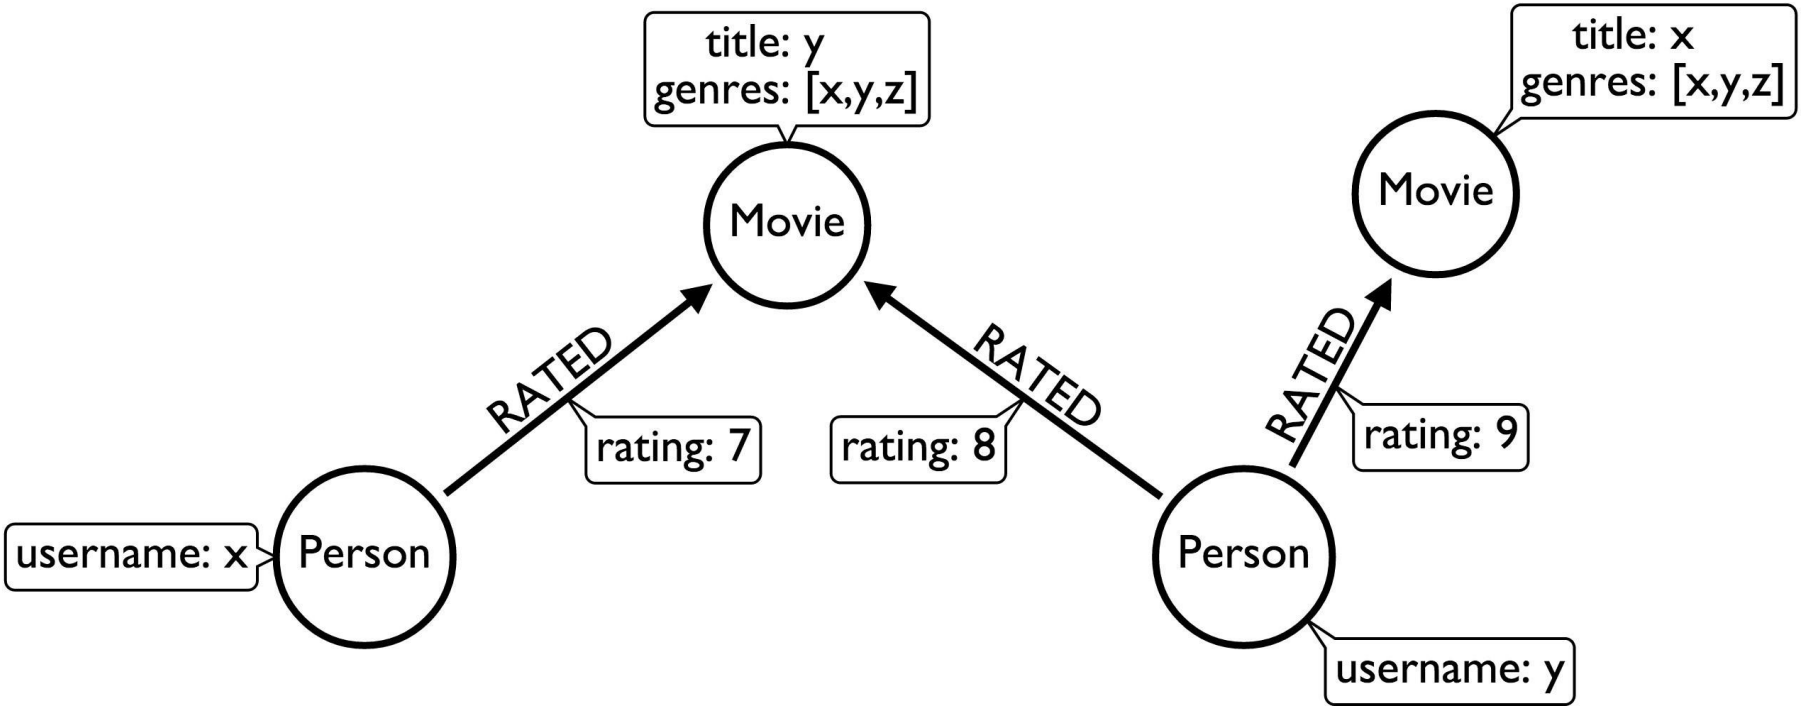

Cypher Query: Movie Recommendation

What are the Top 25 Movies

- that I haven't seen
- with the same genres as Toy Story
- given high ratings
- by women under 35 who liked Toy Story

```
MATCH (watched:Movie {title:"Toy Story”}) <-[r1:RATED]- (p2) -[r2:RATED]-> (unseen:Movie)
WHERE r1.rating > 7 AND r2.rating > 7 AND p2.gender = “female” AND p2.age < 35
AND watched.genres = unseen.genres
AND NOT( (p:Person) -[:RATED|WATCHED]-> (unseen) )
AND p.username in [“maxdemarzi”,”janedoe”,”jamesdean”]
RETURN unseen.title, COUNT(*)
ORDER BY COUNT(*) DESC
LIMIT 25
```

Cypher Query: Ratings of Two Users

What are the Movies these 2 users have both rated

```
MATCH  (p1:Person {name:'Michael Sherman'}) -[r1:RATED]-> (m:Movie),
       (p2:Person {name:'Michael Hunger'}) -[r2:RATED]-> (m:Movie)
RETURN m.name AS Movie,
       r1.rating AS `M. Sherman's Rating`,
       r2.rating AS `M. Hunger's Rating`
```

Movie Data Model

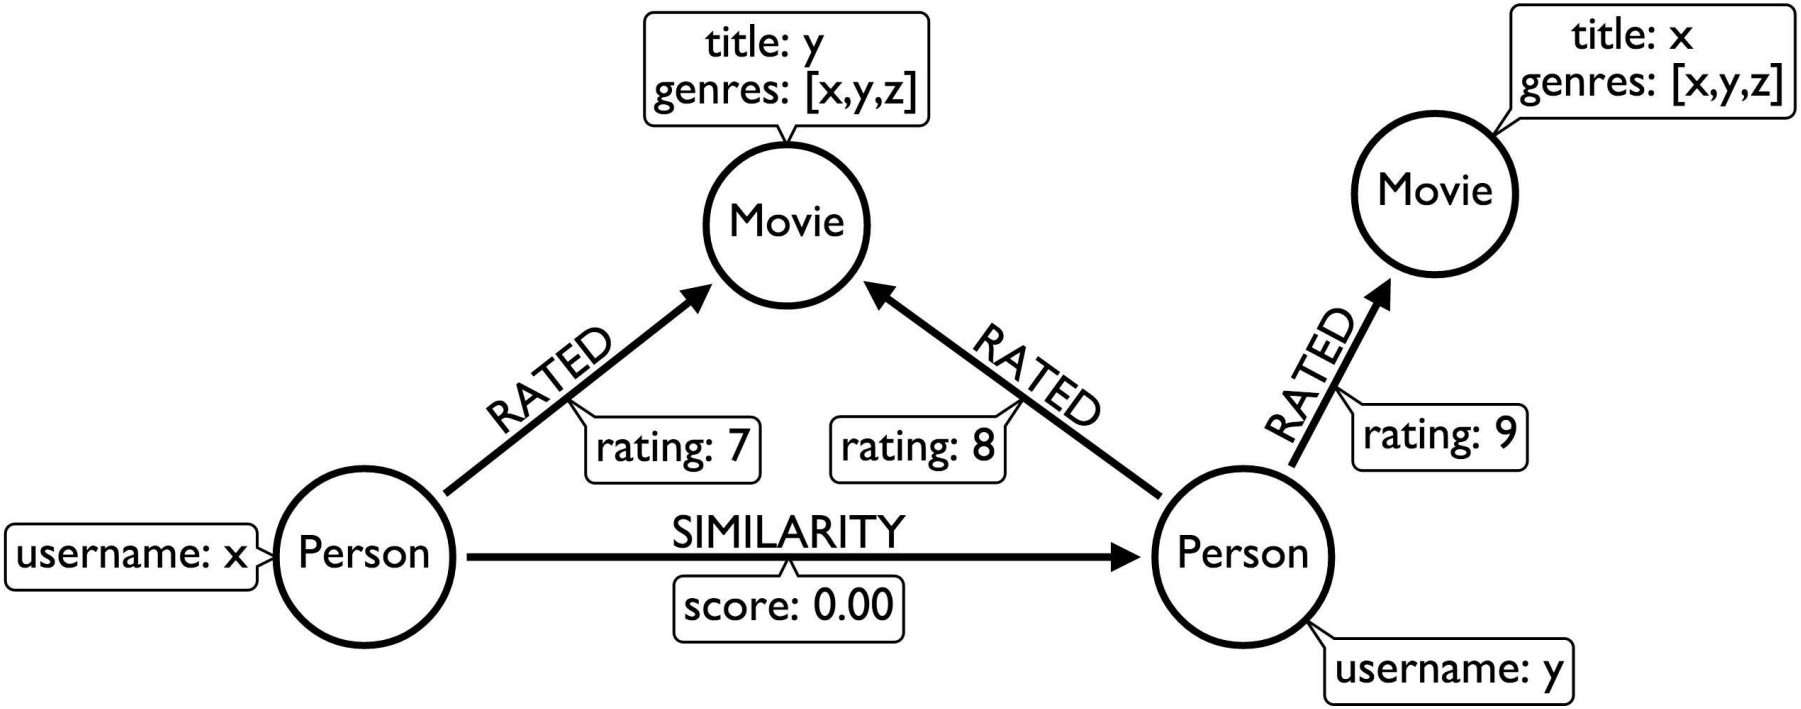

Knowledge Graphs

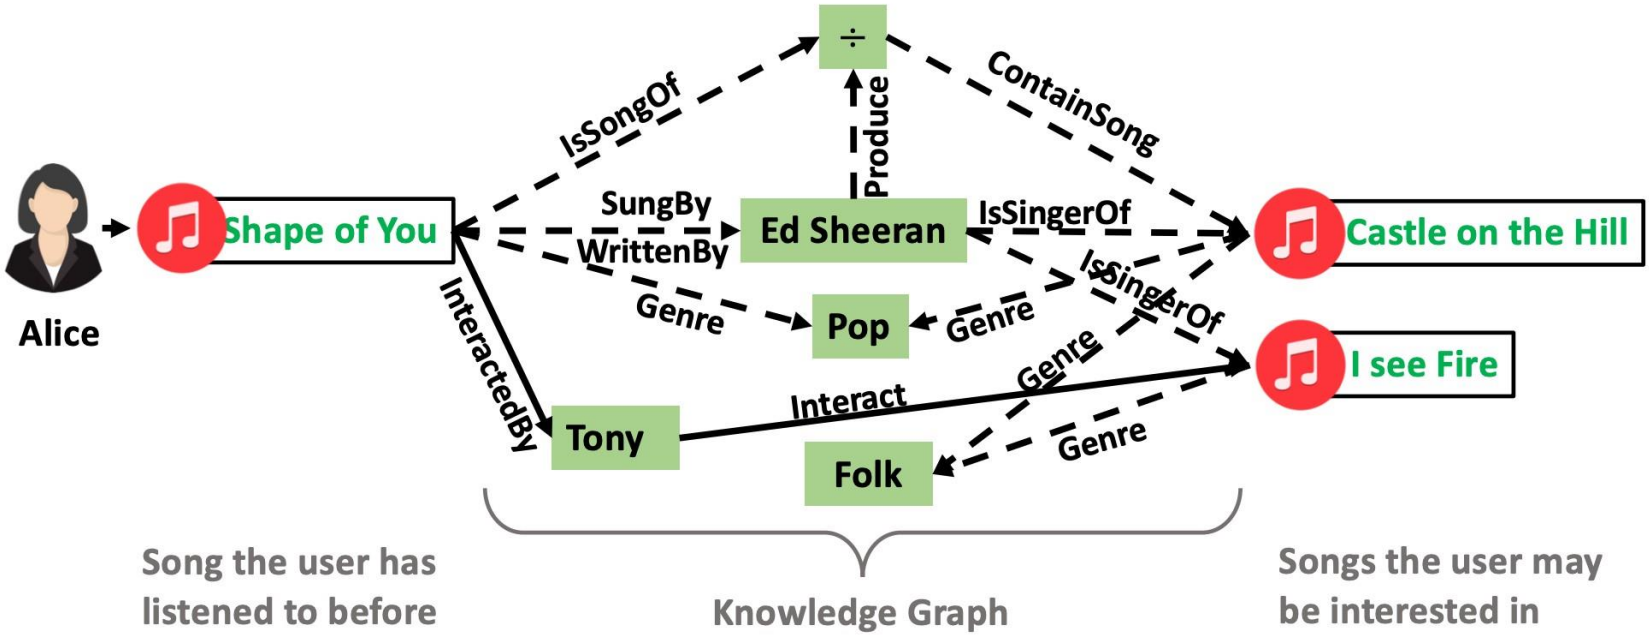

Explainable Reasoning over Knowledge Graphs for Recommendation

Knowledge Graphs

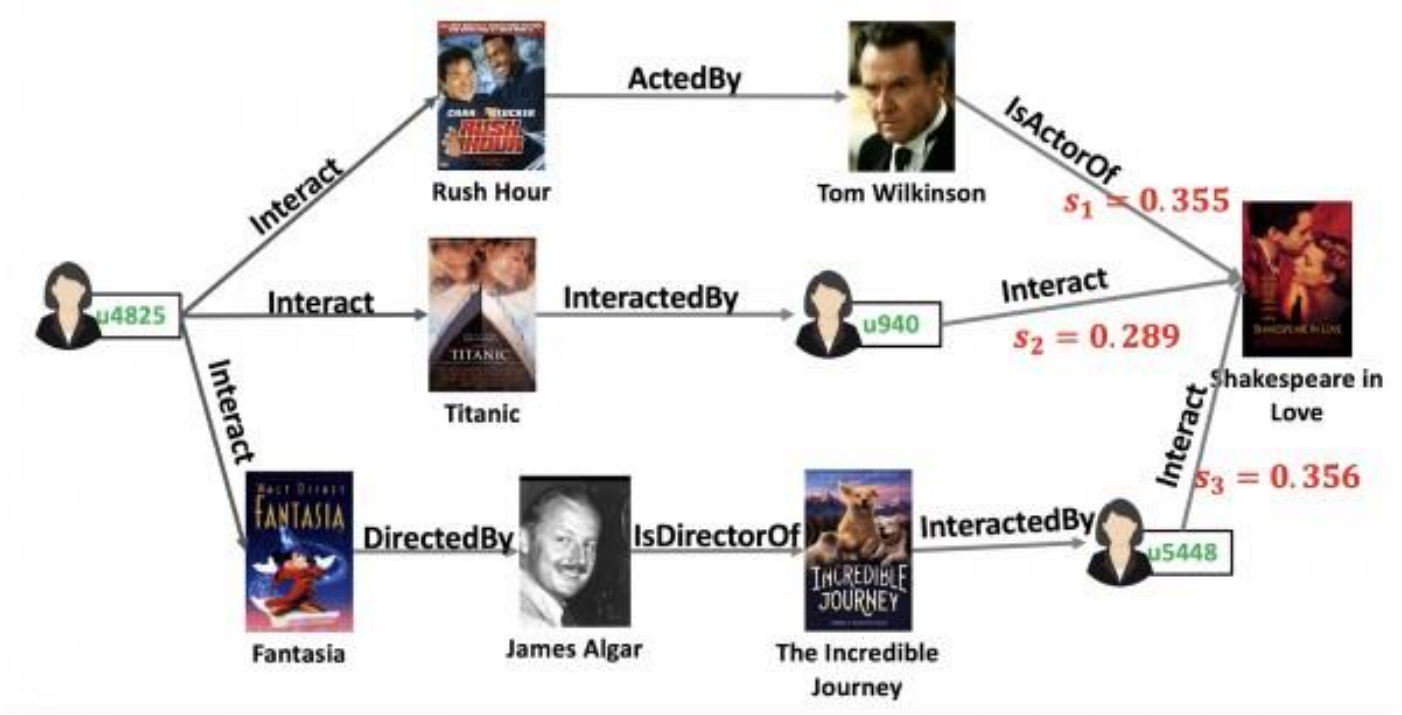

Explainable Reasoning over Knowledge Graphs for Recommendation

Neo4j Developer Page

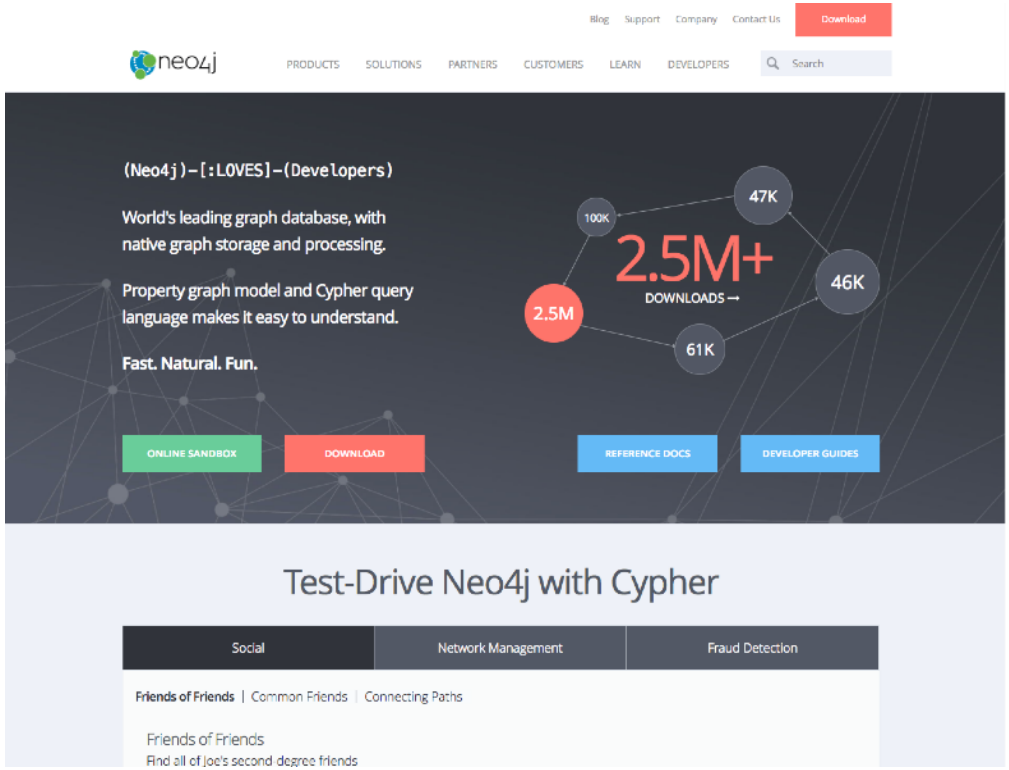

neo4j.com/developer

https://neo4j.com/blog/graph-of-thrones/

Neo4j Sandbox

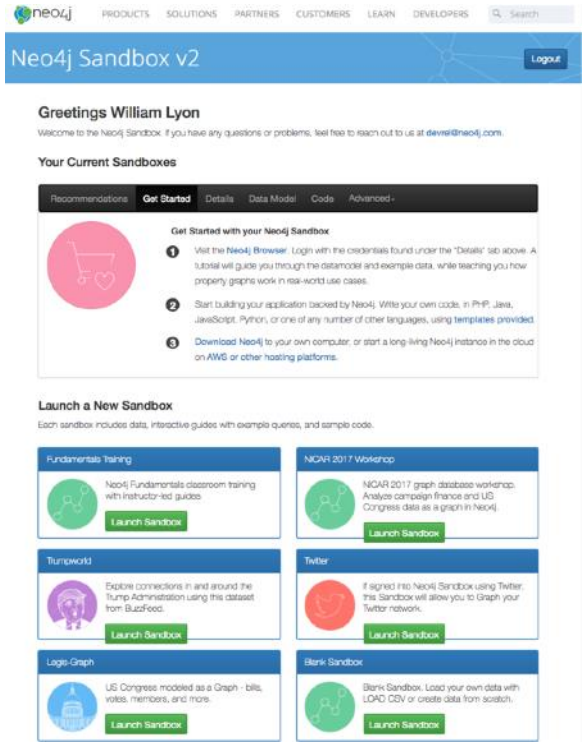

neo4jsandbox.com

Graph Machine Learning

- Machine learning on graphs is an important and ubiquitous task with applications ranging from drug design.
- The primary challenge in this domain is **finding a way to represent, or encode, graph structure so that it can be easily exploited by machine learning models**.
- https://ai.googleblog.com/2016/10/graph-powered-machine-learning-at-google.html
- GNNs (Graph Neural Networks)

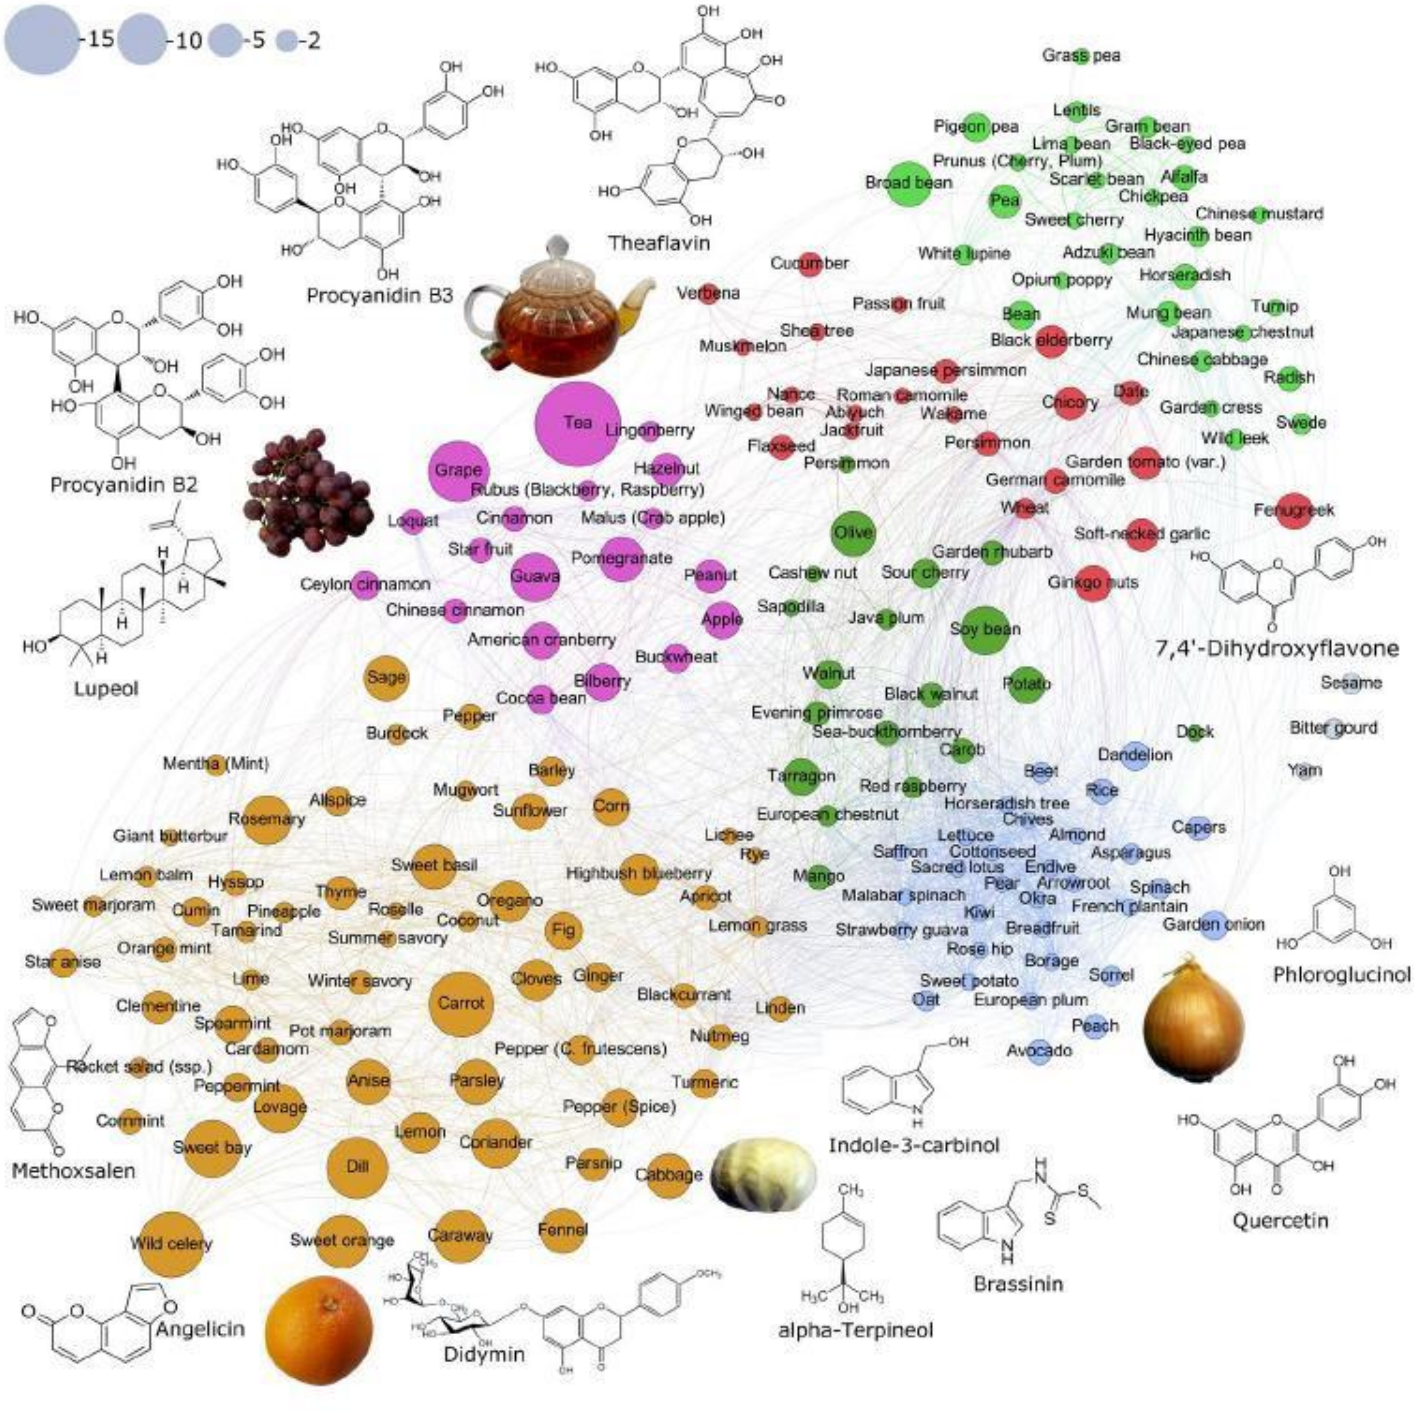
15
10
5
2
Theaflavin
Procyanidin B3
Procyanidin B2
Lupeol
Methoxsalen
Angelicin
Didymin
Indole-3-carbinol
alpha-Terpineol
Brassinin
Quercetin
7,4'-Dihydroxyflavone
Phloroglucinol
Tea
Grape
Rubus (Blackberry, Raspberry)
Hazelnut
Lingonberry
Cinnamon
Malus (Crab apple)
Loquat
Star fruit
Pomegranate
Guava
Ceylon cinnamon
Chinese cinnamon
Peanut
Apple
American cranberry
Buckwheat
Bilberry
Cocoa bean
Pepper
Burdock
Sage
Olive
Cashew nut
Sour cherry
Garden rhubarb
Ginkgo nuts
Sapodilla
Java plum
Soy bean
Walnut
Black walnut
Potato
Evening primrose
Sea-buckthornberry
Carob
Tarragon
Red raspberry
Horseradish tree
European chestnut
Chives
Lettuce
Almond
Capers
Saffron
Cottonseed
Asparagus
Sacred lotus
Endive
Pear
Arrowroot
Spinach
Malabar spinach
Okra
French plantain
Kiwi
Strawberry guava
Breadfruit
Garden onion
Rose hip
Borage
Sorrel
Sweet potato
Oat
European plum
Peach
Avocado
Lichee
Rye
Mango
Apricot
Lemon grass
Highbush blueberry
Sweet basil
Oregano
Coconut
Fig
Summer savory
Roselle
Thyme
Hyssop
Lemon balm
Sweet marjoram
Cumin
Pineapple
Tamarind
Orange mint
Star anise
Lime
Winter savory
Cloves
Ginger
Blackcurrant
Linden
Nutmeg
Pepper (C. frutescens)
Turmeric
Pepper (Spice)
Parsley
Anise
Lovage
Peppermint
Corrmint
Sweet bay
Lemon
Coriander
Parsnip
Cabbage
Dill
Carrot
Clementine
Spearmint
Pot marjoram
Cardamom
Rocket salad (ssp.)
Wild celery
Sweet orange
Caraway
Fennel
Mentha (Mint)
Barley
Mugwort
Sunflower
Corn
Allspice
Rosemary
Giant butterbur
Grass pea
Lentils
Pigeon pea
Gram bean
Lima bean
Black-eyed pea
Prunus (Cherry, Plum)
Broad bean
Scarlet bean
Alfalfa
Chickpea
Pea
Sweet cherry
Chinese mustard
Hyacinth bean
White lupine
Adzuki bean
Cucumber
Opium poppy
Horseradish
Verbena
Passion fruit
Bean
Mung bean
Turnip
Shea tree
Japanese chestnut
Muskmelon
Black elderberry
Chinese cabbage
Japanese persimmon
Radish
Nance
Roman camomile
Chicory
Date
Winged bean
Abyssinian
Wakame
Garden cress
Jackfruit
Swede
Flaxseed
Persimmon
Wild leek
Persimmon
Garden tomato (var.)
German camomile
Wheat
Fenugreek
Soft-necked garlic
Bitter gourd
Dock
Dandelion
Yam
Beet
Rice
Sesame