Data base systems homework 9

1. Create a cluster named 'redshift_hw9' (only use free tier)

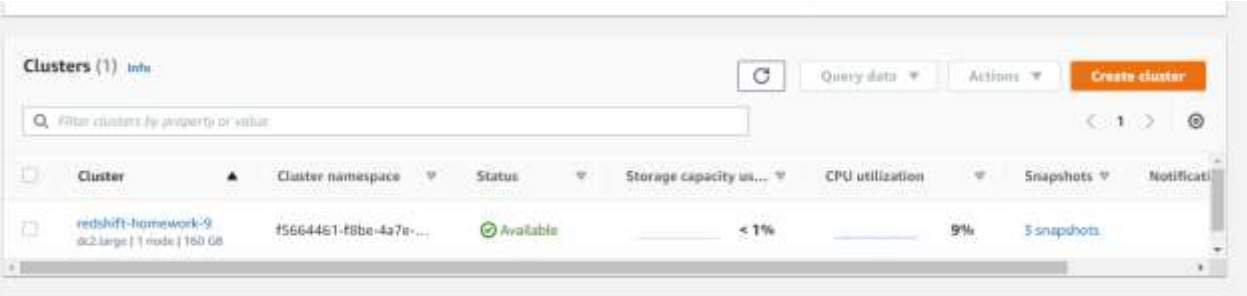

2. Create/Show the required IAM roles for the task.

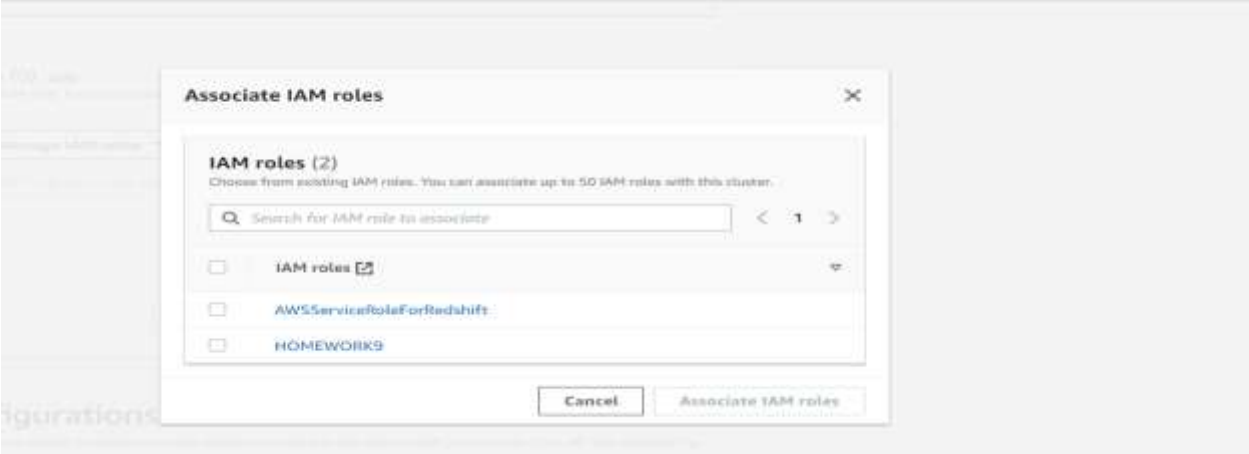

3Create an S3 bucket and load the tickitdb files.

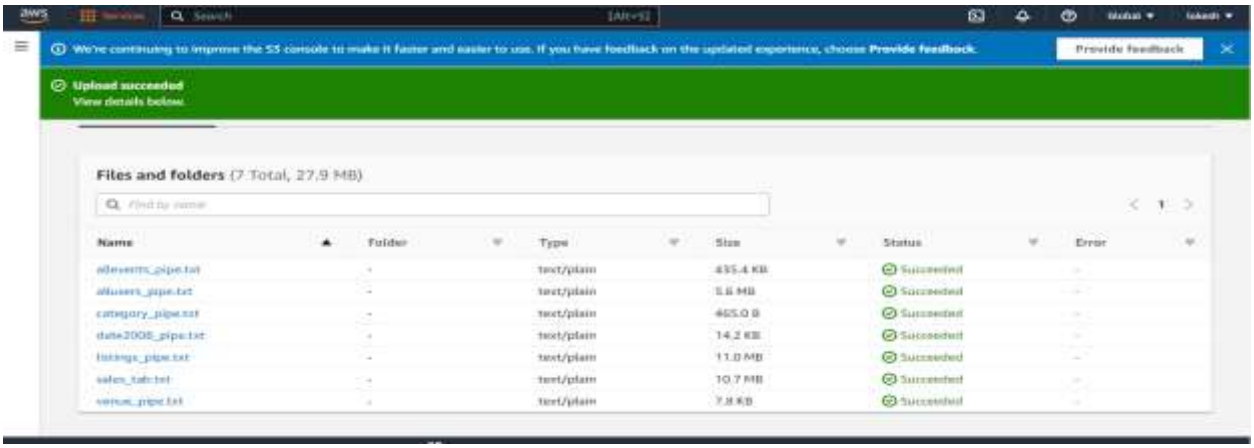

4.Create tables on your cluster. (CATEGORY table, DATE table, EVENT table, VENUE table, USERS table, LISTING table, SALES table)


```
create table homework9.CATEGORY(
CN numeric(10),
categeory varchar(100),
name varchar(100),
league varchar(100)
);
```

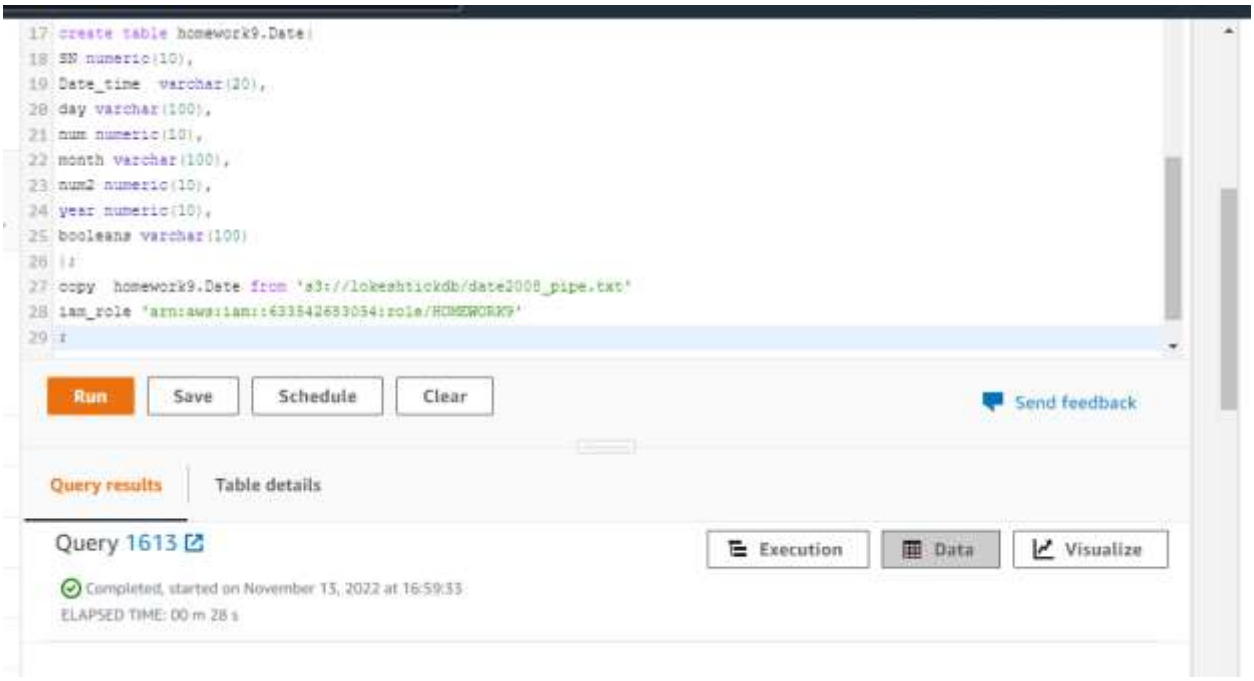

5. Copy the data from S3 to Redshift cluster 'redshift_hw9'.

```
10
11 copy  homework9.CATEGORY from 's3://lokeshtickdb/category_pipe.txt'
12 iam_role 'arn:aws:iam::633542683054:role/HOMEWORK9'
13 ;
```

Rows returned (11) Export

Search rows < 1 2 >

cn	category	name	league
1	Sports	MLB	Major League Baseball
2	Sports	NHL	National Hockey League
3	Sports	NFL	National Football League
4	Sports	NBA	National Basketball Association
5	Sports	MLS	Major League Soccer
6	Shows	Musicals	Musical theatre
7	Shows	Plays	All non-musical theatre
8	Shows	Opera	All opera and light opera
9	Concerts	Pop	All rock and pop music concerts

6.Write a query in the Redshift query editor to find the top 3 sellers in **San Diego, New York, Boston, and San Jose** based on the number of tickets sold in 2008. (Use Window functions)

```
152 SELECT firstname, lastname, total_quantity, venuecity
153 FROM (SELECT s.sellerid, sum(qtysold) total_quantity, v.venuecity,
154  ROW_NUMBER() OVER(PARTITION BY total_quantity ORDER BY total_quantity DESC) AS
155 row_num
156  FROM sales s
157  JOIN event e ON s.eventid = e.eventid
158  JOIN venue v ON v.venueid = e.venueid
159  GROUP BY s.sellerid, v.venuecity
160  HAVING (v.venuecity = 'New York City')
161  ORDER BY total_quantity desc limit 3) Q, users
162 WHERE Q.sellerid = userid
163 ORDER BY Q.total_quantity desc;
164
165 SELECT firstname, lastname, total_quantity, venuecity
166 FROM (SELECT s.sellerid, sum(qtysold) total_quantity, v.venuecity,
167  ROW_NUMBER() OVER(PARTITION BY total_quantity ORDER BY total_quantity DESC) AS
168 row_num
169  FROM sales s
170  JOIN event e ON s.eventid = e.eventid
171  JOIN venue v ON v.venueid = e.venueid
172  GROUP BY s.sellerid, v.venuecity
173  HAVING (v.venuecity = 'Boston')
174  ORDER BY total_quantity desc limit 3) Q, users
175 WHERE Q.sellerid = userid
176 ORDER BY Q.total_quantity desc;
177
178 SELECT firstname, lastname, total_quantity, venuecity
```

Data base systems homework 9

```
177
178 SELECT firstname, lastname, total_quantity, venuecity
179 FROM (SELECT s.sellerid, sum(qtysold) total_quantity, v.venuecity,
180  ROW_NUMBER() OVER(PARTITION BY total_quantity ORDER BY total_quantity DESC) AS
181 row_num
182  FROM sales s
183  JOIN event e ON s.eventid = e.eventid
184  JOIN venue v ON v.venueid = e.venueid
185  GROUP BY s.sellerid, v.venuecity
186  HAVING (v.venuecity = 'San Jose')
187  ORDER BY total_quantity desc limit 3) Q, users
188 WHERE Q.sellerid = userid
189 ORDER BY Q.total_quantity desc;
190
191
```

Run Save Schedule Clear Send feedback

Query results | Table details

Query 21639

Completed, started on November 14, 2022 at 23:18:49

Execution Data Visualize

Rows returned (3)

Export

Search rows 1

firstname	lastname	total_quantity	venuecity
Yeo	Gregory	13	San Diego
Sylvester	Rutledge	12	San Diego
Yoshio	Lester	11	San Diego

Search rows 1

firstname	lastname	total_quantity	venuecity
Logan	David	27	New York City
Patricia	Espinoza	25	New York City
Summer	Marsh	25	New York City

firstname	lastname	total_quantity	venuecity
Teagan	Delaney	17	Boston
Kessie	Hunter	17	Boston
Owen	Hyde	15	Boston
Hilda	Dudley	17	San Jose
Hedy	Dickerson	13	San Jose
Amethyst	Morgan	12	San Jose

7) Write a query in the Redshift query editor to see events with the lowest sales.

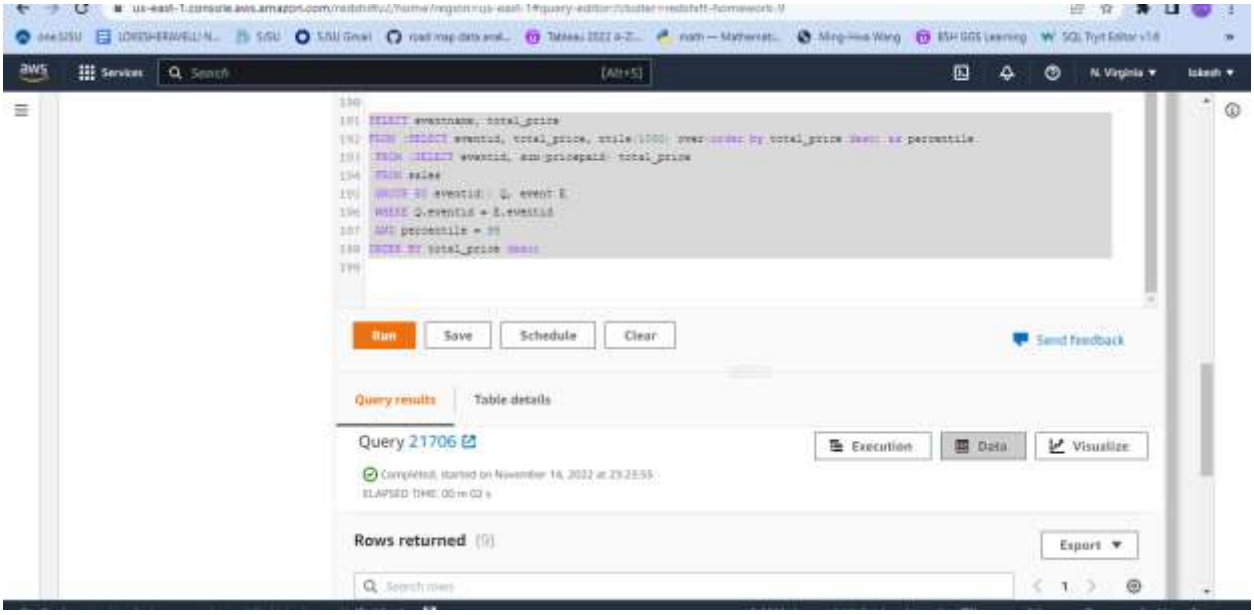

Query 21706

Completed, started on November 14, 2022 at 23:23:55

ELAPSED TIME: 00 m 02 s

Execution | Data | Visualize

Rows returned (9)

Export

Search rows

1

eventname	total_price
Rock 'n' Roll	22356.00
Ben Folds	22342.00
Endgame	22342.00
Bow Wow	22334.00
A Bronx Tale	22334.00
Commodores	22331.00
Blue Man Group	22327.00

8) Write a query in the Redshift query editor to count the number of users in each state and city combination.

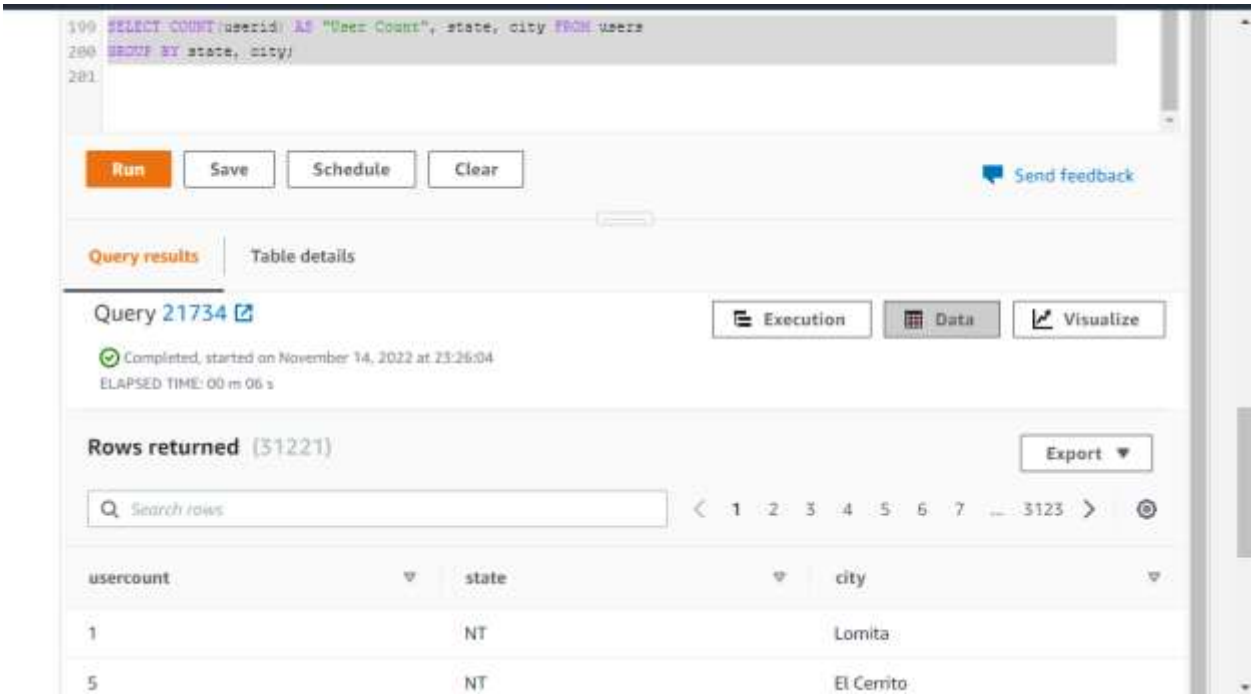

usercount	state	city
1	NT	Lomita
5	NT	El Cerrito
3	NU	Henderson
2	PE	Garland
2	NC	Pueblo
2	MB	Little Rock
2	AB	Hidden Hills
1	MO	Williamsport
2	MB	Waukegan
2	IN	Lebanon

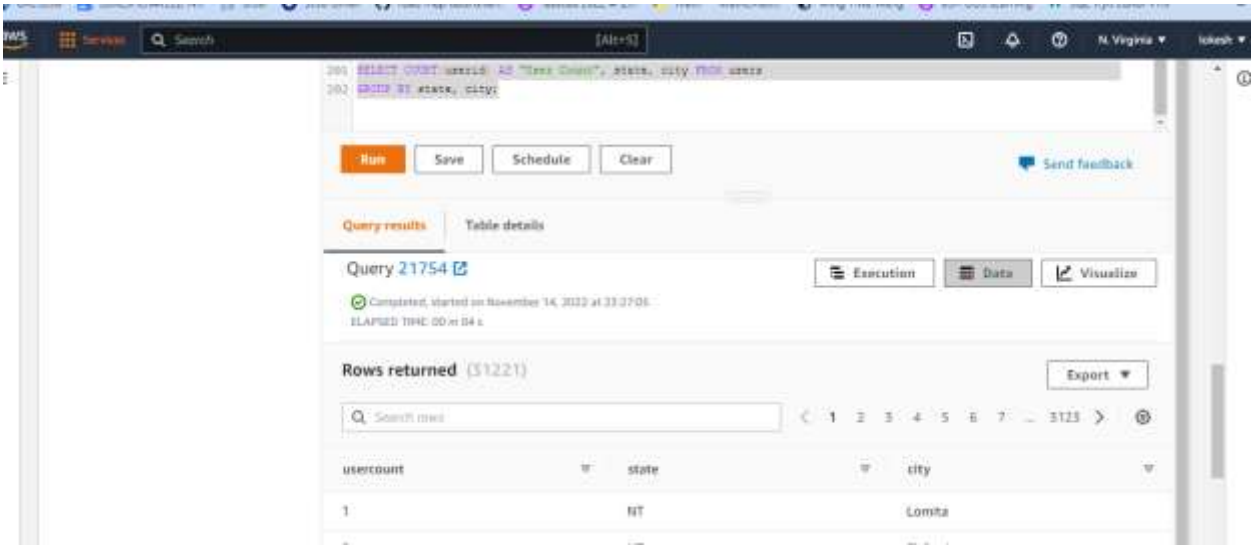

9) What are Database Querying Options available in Amazon Redshift?

10) Delete the cluster 'redshift_hw9', IAM role and S3 bucket

Data base systems homework 9

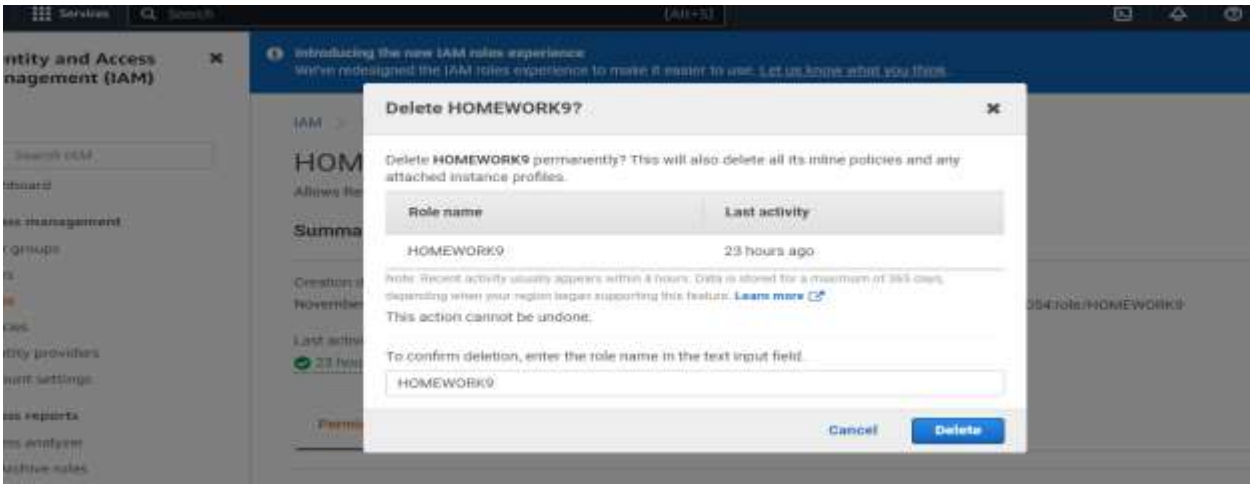

Data base systems homework 9